

Group29_Final

May 13, 2026

Data Mining 2026 — Project Notebook

Course: Data Mining

Project track: Track A: Standard Analysis Project

Group members:

- Victor Bjerggaard Thøstesen
- Nikolai Mose Sannerum
- Frederik Hansen Baattrup

Dataset:

During the project we look at game statistics for Premier league matches since 2005. The datasets from 2003/2004 and 2004/2005 seasons are disregarded due to reading issues.

Some features of the full datasets are combined into a master dataframe that is used in part I and part II. This dataframe contains the features: “HomeTeam”, “AwayTeam”, “Full Time Result (FTR)”, “Referee”, “Full Time Home Goal (FTHG)”, “Full Time Away Goal (FTAG)”, “Home Shots (HS)”, “Away Shots (AS)”, “Home Shots On Target (HST)”, “Away Shots On Target (AST)”, “Home Fouls (HF)”, “Away Fouls (AF)”, “Home Corner (HC)”, “Away Corner (AC)”, “Home Yellow (HY)”, “Away Yellow (Ay)”, “Home Red (HR)”, “Away Red (AR)”, “Average Home Win Odds (AvgH)”, “Average Draw Odds (AvgD)”, “Average Away Win Odds (AvgA)”, “Date” and “Season”.

For part III more features were needed to conduct the pattern mining analysis, so the data was reloaded to extract additional bookmaker odds features.

Dataset Overview: - Source: https://www.football-data.co.uk/englandm.php?fbclid=IwY2xjawQVBq9leHRuAQTbB8sBIRDk8vETIRUAA3XgndZg_aem_OCN-0ddvAe0Fob_1GlxMdg

- Number of instances (matches): 7860
- Number of features / entities: 23 plus additional odds features for part III
- Missing values: No missing values after preprocessing.

0. Reproducibility and Setup

```
[1]: import sys
      print(sys.version)

      import numpy as np
      import pandas as pd
```

```

import os
import math
import matplotlib.pyplot as plt
import networkx as nx
from matplotlib.patches import Patch
from tqdm import tqdm
import scipy.sparse as sp
from sklearn.neighbors import kneighbors_graph
from sklearn.cluster import DBSCAN, KMeans
from sklearn.manifold import TSNE
from sklearn.metrics import adjusted_rand_score, normalized_mutual_info_score
from scipy.spatial.distance import cdist
from itertools import combinations, chain
from joblib import Parallel, delayed

```

```

RANDOM_SEED = 42
np.random.seed(RANDOM_SEED)

```

3.13.7 (v3.13.7:bcee1c32211, Aug 14 2025, 19:10:51) [Clang 16.0.0
(clang-1600.0.26.6)]

1 Clustering (Part I)

Clustering objective:

For the first part of the project we want to analyze our dataset using clustering. Our tentative analysis goal is to pick out a set of features from the dataset which can be meaningfully clustered and compared to a ground truth. First we apply unsupervised clustering to group football matches using in game statistics. The chosen features will be combined and transformed in order to reduce dimensions and make the matches more comparable. Outcome-related variables FTHG, FTAG, FTR, HTR and odds are excluded from clustering to prevent data leak and reserved for evaluation. After normalizing the numeric features, clustering algorithms are applied to identify match groupings. The FTR column serves as ground truth to assess alignment between clusters and actual outcomes (home win, away win, draw). Cluster evaluation metrics includes contingency tables and cluster purity measurements.

1.1 Dataset Description and Loading

```

[2]: # --- Data Extraction Function ---
# Here we define a function to extract the chosen relevant data features from
# raw CSV files 1 for each season.
# This function will look for the average odds columns and handle cases where
# they might be missing.

def extract_data(df, filename):
    """

```

*Extracts premier league data from CSV file.
 Prefers BbAvH/D/A, otherwise falls back to AvgH/D/A.
 Adds season column extracted from filename.*

parameters:

df: DataFrame containing the raw data from the CSV file.

filename: The name of the CSV file, used to extract the season information.

return: A DataFrame with the extracted features and season information, or `None` if required columns are missing.

```

"""

primary_cols = ['HomeTeam', 'AwayTeam', 'FTR', 'Referee', 'FTHG',
                'FTAG', 'HS', 'AS', 'HST', 'AST', 'HF', 'AF',
                'HC', 'AC', 'HY', 'AY', 'HR', 'AR', "BbAvH", "BbAvD", "BbAvA",
↪ 'Date']
fallback_cols = ['HomeTeam', 'AwayTeam', 'FTR', 'Referee', 'FTHG',
                 'FTAG', 'HS', 'AS', 'HST', 'AST', 'HF', 'AF',
                 'HC', 'AC', 'HY', 'AY', 'HR', 'AR', "AvgH", "AvgD", "AvgA",
↪ 'Date']

season = os.path.splitext(filename)[0]

# Trying to find Betbrain columns first
if all(col in df.columns for col in primary_cols):
    result = df[primary_cols].copy()
    result.columns = ['HomeTeam', 'AwayTeam', 'FTR', 'Referee', 'FTHG',
                     'FTAG', 'HS', 'AS', 'HST', 'AST', 'HF', 'AF',
                     'HC', 'AC', 'HY', 'AY', 'HR', 'AR', "AvgH", "AvgD",
↪ "AvgA", 'Date']
    print(f'Found primary columns: BbAvH, BbAvD, BbAvA for {filename}')

# Else trying to find fallback columns
elif all(col in df.columns for col in fallback_cols):
    result = df[fallback_cols].copy()
    print(f'Using fallback columns: AvgH, AvgD, AvgA for {filename}')

# If neither set of columns is found, print a warning and return None
else:
    print(f" Warning: Required odds columns not found in {filename}")
    return None

# Add season column
result["Season"] = season

return result

```

```
[3]: # --- User settings begin ---
dataPath = "Data" # Define the folder where the raw CSV files are located.
# --- User settings end ---

new_data = []

for file in os.listdir(dataPath):
    if file.endswith(".csv"):
        path = os.path.join(dataPath, file)
        try:
            df = pd.read_csv(path)
        except Exception as e:
            print(f"Error reading {file}: {e}")
            continue

        extracted = extract_data(df, file)

        if extracted is not None:
            new_data.append(extracted)

# Combine everything into one dataframe
final_df = pd.concat(new_data, ignore_index=True)
```

```
Found primary columns: BbAvH, BbAvD, BbAvA for 20072008premierleague.csv
Found primary columns: BbAvH, BbAvD, BbAvA for 20112012premierleague.csv
Found primary columns: BbAvH, BbAvD, BbAvA for 20132014premierleague.csv
Using fallback columns: AvgH, AvgD, AvgA for 20242025premierleague.csv
Error reading 20042005premierleague.csv: 'utf-8' codec can't decode byte 0xa0 in
position 75614: invalid start byte
Using fallback columns: AvgH, AvgD, AvgA for 20202021premierleague.csv
Using fallback columns: AvgH, AvgD, AvgA for 20192020premierleague.csv
Found primary columns: BbAvH, BbAvD, BbAvA for 20082009premierleague.csv
Found primary columns: BbAvH, BbAvD, BbAvA for 20062007premierleague.csv
Found primary columns: BbAvH, BbAvD, BbAvA for 20152016premierleague.csv
Using fallback columns: AvgH, AvgD, AvgA for 20222023premierleague.csv
Found primary columns: BbAvH, BbAvD, BbAvA for 20102011premierleague.csv
Found primary columns: BbAvH, BbAvD, BbAvA for 20092010premierleague.csv
Found primary columns: BbAvH, BbAvD, BbAvA for 20122013premierleague.csv
Using fallback columns: AvgH, AvgD, AvgA for 20252026premierleague.csv
Error reading 20032004premierleague.csv: Error tokenizing data. C error:
Expected 57 fields in line 305, saw 72
```

```
Found primary columns: BbAvH, BbAvD, BbAvA for 20172018premierleague.csv
Found primary columns: BbAvH, BbAvD, BbAvA for 20162017premierleague.csv
Found primary columns: BbAvH, BbAvD, BbAvA for 20052006premierleague.csv
Using fallback columns: AvgH, AvgD, AvgA for 20232024premierleague.csv
Using fallback columns: AvgH, AvgD, AvgA for 20212022premierleague.csv
Found primary columns: BbAvH, BbAvD, BbAvA for 20182019premierleague.csv
```

Found primary columns: BbAvH, BbAvD, BbAvA for 20142015premierleague.csv

```
[4]: print(f"\n Final dataset:")
final_df = final_df.dropna(subset=["HomeTeam", "AwayTeam"]) # 2014/2015 season
    ↪has a row of NaN that is removed here
final_df
```

Final dataset:

```
[4]:      HomeTeam    AwayTeam FTR      Referee FTHG FTAG   HS   AS \
0      Aston Villa    Liverpool   A      M Riley   1.0   2.0  10.0  17.0
1          Bolton    Newcastle   A        C Foy   1.0   3.0  13.0   7.0
2          Derby    Portsmouth   D        M Dean   2.0   2.0  12.0  12.0
3      Everton      Wigan      H M Clattenburg   2.0   1.0  12.0  14.0
4  Middlesbrough    Blackburn   A    A Marriner   1.0   2.0  10.0   4.0
...
7855      Hull    Man United   D    L Probert   0.0   0.0  16.0   7.0
7856    Leicester      QPR      H    M Oliver   5.0   1.0  22.0  18.0
7857    Man City    Southampton   H        C Foy   2.0   0.0  15.0  13.0
7858    Newcastle    West Ham   H    M Atkinson   2.0   0.0  17.0   4.0
7859      Stoke    Liverpool   H    A Taylor   6.0   1.0  15.0  13.0
```

```
      HST  AST  ...  AC  HY  AY  HR  AR  AvgH  AvgD  AvgA      Date \
0      6.0  7.0  ...  2.0  4.0  2.0  0.0  0.0  3.67  3.23  1.95  11/08/07
1      9.0  5.0  ...  3.0  1.0  1.0  0.0  0.0  2.41  3.18  2.80  11/08/07
2      5.0  6.0  ...  6.0  1.0  2.0  0.0  0.0  2.85  3.19  2.38  11/08/07
3      8.0  4.0  ...  2.0  0.0  0.0  0.0  0.0  1.66  3.45  5.25  11/08/07
4      6.0  4.0  ...  3.0  3.0  4.0  0.0  0.0  2.41  3.18  2.81  11/08/07
...
7855  6.0  1.0  ...  1.0  2.0  2.0  0.0  1.0  3.14  3.55  2.21  24/05/15
7856  7.0  2.0  ...  6.0  0.0  0.0  0.0  0.0  1.56  4.21  5.71  24/05/15
7857  6.0  4.0  ...  4.0  1.0  1.0  0.0  0.0  1.56  4.27  5.61  24/05/15
7858  4.0  1.0  ...  3.0  2.0  1.0  0.0  0.0  1.77  3.84  4.42  24/05/15
7859  9.0  4.0  ...  9.0  4.0  2.0  0.0  0.0  3.42  3.48  2.12  24/05/15
```

```
      Season
0  20072008premierleague
1  20072008premierleague
2  20072008premierleague
3  20072008premierleague
4  20072008premierleague
...
7855  20142015premierleague
7856  20142015premierleague
7857  20142015premierleague
7858  20142015premierleague
7859  20142015premierleague
```

[7860 rows x 23 columns]

1.2 Ground truth encoding

```
[5]: def encode_results(results):  
    """  
    Endcoding of results  
    1: HOME WIN, -1: AWAY WIN, 0: DRAW  
  
    Parameters:  
        results (pandas Series): results to encode  
  
    Returns:  
        encoding (numpy array): encoded results  
    """  
    encoding = np.zeros(len(results)) # Initialize encoding array  
  
    for i in range(len(results)):  
        if results[i] == "H":  
            encoding[i] = 1 # Home win  
        elif results[i] == "A":  
            encoding[i] = -1 # Away win  
        else:  
            encoding[i] = 0 # Draw  
    return encoding  
data = final_df.copy()  
ground_truth = encode_results(data['FTR'])  
print("Amount of home wins: ", np.sum(ground_truth == 1))  
print("Amount of away wins: ", np.sum(ground_truth == -1))  
print("Amount of draws: ", np.sum(ground_truth == 0))
```

Amount of home wins: 3586

Amount of away wins: 2368

Amount of draws: 1906

1.3 Data preprocessing and feature engineering

1.3.1 New Features and Normalization

To make matches more comparable, new features were engineered as the difference between home and away statistics. A positive value indicates that the home team has a higher value for the given metric.

The following features were constructed:

- Shot difference ($HS - AS$)
- Shots on target difference ($HST - AST$)
- Foul difference ($HF - AF$)

- Corner difference ($HC - AC$)
- Difference in average betting odds ($AvgH - AvgA$)

After feature engineering, the variables were **normalized using z-score normalization** to ensure comparable scales across features.

1.3.2 Data insights:

Finally, histograms of the normalized features were plotted to inspect their distributions. The following insights are observed: - All variables appear roughly symmetric and centered around zero, which is expected after z-score normalization. This indicates that, on average, neither the home nor away team consistently dominates these metrics across matches.

- Most feature distributions appear approximately Gaussian after normalization. The exceptions are RC_diff and YC_diff, which are highly discrete due to the low and integer-valued nature of card events. The largely unimodal and continuous shape of the other features suggests that the data may not naturally separate into distinct groups across all variables. Any potential clustering structure may therefore rely more heavily on the discrete features.


Overall, this indicates that clusters, if present, are likely to be weakly separated, which may make clustering more challenging.

```
[6]: # --- Creating new features by calculating the difference between home and away
      ↪statistics (Positive means home team have more) ---
modified_stats = pd.DataFrame({'SD': data['HS'] - data['AS'],
                              'SOTD': data['HST'] - data['AST'],
                              'F_diff': data['HF'] - data['AF'],
                              'C_diff': data['HC'] - data['AC'],
                              'YC_diff': data['HY'] - data['AY'],
                              'AvgOdds_diff': data['AvgH'] - data['AvgA'],
                              'RC_diff': data['HR'] - data['AR']
                              })

# --- Normalize the new features with average and standard deviation (Z-score
      ↪normalization) ---
modified_stats = (modified_stats - modified_stats.mean()) / modified_stats.std()

# -- Plot as histogram to see the distribution of the new features ---
modified_stats.hist(bins=30, figsize=(12, 8))
plt.suptitle('Distribution of Modified Features (Normalized)', fontsize=16)
plt.show()
```

Distribution of Modified Features (Normalized)



1.3.3 Scatter Plots:

The scatter plots below show the relationships between the normalized features and the ground truth, Full Time Result (FTR), which indicates whether the home team won (1), drew (0), or lost (-1). Each scatter plot illustrates how various match statistics, such as shot differences, foul differences, and card differences, are distributed against FTR:

- **SD vs FTR**: Shot difference vs Full Time Result
- **SOTD vs FTR**: Shots on target difference vs Full Time Result
- **F_diff vs FTR**: Foul difference vs Full Time Result
- **C_diff vs FTR**: Corner difference vs Full Time Result
- **AvgOdds_diff vs FTR**: Difference in average betting odds vs Full Time Result

Each scatter plot demonstrates the relationship between these engineered features and FTR (ground truth). This was done to get an idea of which features that seems important when comparing to ground truth. It seems that there are some correlation between some of the features and ground truth, though they still might be hard to cluster.


```
[7]: n_cols = 3 # number of columns in the subplot grid
     n_features = len(modified_stats.columns)
     n_rows = math.ceil(n_features / n_cols)
```



```
plt.figure(figsize=(4 * n_cols, 3 * n_rows))

for i, column in enumerate(modified_stats.columns):
    plt.subplot(n_rows, n_cols, i + 1)
    plt.scatter(modified_stats[column], ground_truth, alpha=0.1)
    plt.title(f'{column} vs FTR')
    plt.xlabel(column)
    plt.ylabel('FTR')

plt.tight_layout()
plt.show()
```



1.4 DBSCAN

1.4.1 Tuning the parameter ϵ

Before running DBSCAN on the dataset, we try to find the best hyperparameter for ϵ .

First we want to find the “border object” or “first valley” (non differentiable point). This is done by defining a function that for each point find the distance to the k -nearest neighbor. This distance is the distance to the MinPts-nearest neighbor where MinPts is set to $2d - 1$ where d is the number of features (Slide 22 lecture 3).

The distance measure used, is both the euclidean distance and manhattan distance.

Afterwards we plot the distances in descending order.

```
[8]: # -- Calculating K nearest neighbours --
def find_k_nearest_neighbor_dist(data, dist: str = 'euclidean', minpts=None,
    use_tqdm=True):
    """
    Compute the distance to the k-th nearest neighbor for each point.

    Parameters:
        data (numpy array): Shape (n_samples, d_features)
        dist (str): The distance metric to use (default 'euclidean')
        minpts (int): Number of neighbors to consider (default 2*d - 1 -- See
    ↪ slide 22 lecture 3)
        use_tqdm (bool): Show progress bar for large datasets
        minpts (int): Number of neighbors to consider (default 2*d - 1 -- See
    ↪ slide 22 lecture 3)
        use_tqdm (bool): Show progress bar for large datasets

    Returns:
        distances_sorted (numpy array): k-th nearest neighbor distances sorted
    ↪ descending
        minpts (int): The k used
    """
    n, d = data.shape
    if minpts is None:
        minpts = 2 * d - 1

    # Compute full pairwise distance matrix (n x n)
    pairwise_dist = cdist(data, data, metric=dist) # Manhattan distance is used

    np.fill_diagonal(pairwise_dist, np.inf) # Set diagonal to np.inf to ignore
    ↪ self-distance

    distances = np.zeros(n)

    iterator = tqdm(range(n), disable=not use_tqdm)
    for i in iterator:
        # Sort distances for this point
        sorted_dists = np.partition(pairwise_dist[i], minpts) # faster than
    ↪ full sort
        distances[i] = sorted_dists[minpts] # minpts-th nearest neighbor

    # Return distances sorted descending
    return np.sort(distances)[::-1], minpts
```

```

y_man, minpts_man = find_k_nearest_neighbor_dist(modified_stats.values,
    ↪dist='cityblock')
print('Found value for minpts (Manhattan):', minpts_man)

y_euclidean, minpts_euclidean = find_k_nearest_neighbor_dist(modified_stats.
    ↪values, dist='euclidean')
print('Found value for minpts (Euclidean):', minpts_euclidean)

```

100%| | 7860/7860 [00:00<00:00, 17344.33it/s]

Found value for minpts (Manhattan): 13

100%| | 7860/7860 [00:00<00:00, 17850.99it/s]


Found value for minpts (Euclidean): 13

```

[9]: # -- Plotting K nearest Neighbor distances---
x_man = np.arange(len(y_man))
x_euclidean = np.arange(len(y_euclidean))
epsilon_man = 2.7
epsilon_euclidean = 1.8
fig, ax = plt.subplots(1,2, figsize=(10, 5))
ax[0].plot(x_man, y_man)
ax[0].set_xlabel("Point Index")
ax[0].set_ylabel(f"Distance to {minpts_man}th Nearest Neighbor")
ax[0].set_title(f"Manhattan distance")
ax[0].axhline(epsilon_man, c='k', ls='--', label=f'Epsilon = {epsilon_man}')
ax[0].legend()

ax[1].plot(x_euclidean, y_euclidean)
ax[1].set_xlabel("Point Index")
ax[1].set_ylabel(f"Distance to {minpts_euclidean}th Nearest Neighbor")
ax[1].set_title(f"Euclidean distance")
ax[1].axhline(epsilon_euclidean, c='k', ls='--', label=f'Epsilon =
    ↪{epsilon_euclidean}')
ax[1].legend()
plt.tight_layout()

```



As seen above, it is hard to determine the right ϵ as there is no clear valley in either of the plots. This is a sign that there isn't clearly defined clusters in the data, as the density acts continuously rather than discretely. However we choose $\epsilon = 1.8$ for the euclidean distance and $\epsilon = 2.7$ for the manhattan distance, and continue to DBSCAN.

1.4.2 DBSCAN clustering

```
[10]: dbscan = DBSCAN(eps=epsilon_euclidean, min_samples=minpts_euclidean)
dbscan.fit(modified_stats.values)
labels_dbscan = dbscan.labels_
print("DBSCAN Cluster Labels:", np.unique(labels_dbscan)) # -1 corresponds to
↪noise points, other integers correspond to cluster labels
```

DBSCAN Cluster Labels: [-1 0 1 2]

It is observed that the result of DBSCAN when using the manhattan distance is extremely sensitive to the choice of epsilon. Choosing: - $\epsilon = 2.8$ results in 1 cluster and some noise - $\epsilon = 2.75$ results in 2 clusters and some noise. - $\epsilon = 2.7$ results in 4 clusters and some noise.

This is a sign that the data is ill suited for clustering using the manhattan distance. When using the euclidean distance, the result is more stable. Therefore we choose to continue with the euclidean distance for the rest of the analysis.

There is no clear valley in the plot of the k-nearest neighbor distances, so the clusters obtained may not be well defined. As seen above we obtain three clusters using the euclidean distance (0, 1, 2) and some noise (-1).

To evaluate the clustering, we define two functions, one makes the contingency table based on the FTR ground truth, and the other evaluates the purity score. - The contingency table tells us how many points of each ground truth the clusters contain. - The purity measure tells us how many points in the cluster belong to the same (most prominent) ground truth.

```

[11]: def contingency_table(C, T, ignore_noise=True):
    """
    Parameters:
        C(numpy array):      Clusters obtained by a clustering algorithm as a
        ↪ nx1 vector
        T(numpy array):      Ground-truth cluster assignments as a nx1 vector

    Returns:
        ctable (numpy array): a num_clusters_of_C x num_cluster_of_T matrix
        ↪ containing the overlaps among the different clusters
    """

    if ignore_noise:
        mask = (C != -1)
        C_masked = C[mask]
        indices_masked = np.where(mask)[0]
    else:
        C_masked = C
        indices_masked = np.arange(len(C))

    clusters = np.unique(C_masked)
    classes = np.unique(T) # include all classes, not just masked

    ctable = np.zeros((len(clusters), len(classes)), dtype=int)

    for i, c in enumerate(clusters):
        for j, t in enumerate(classes):
            # count overlap only for points that are not noise
            count = np.sum((C_masked == c) & (T[indices_masked] == t))
            ctable[i, j] = count
    return ctable

def purity_score(ctable):
    """
    Parameters:
        ctable (numpy array): Contingency table of shape (num_clusters,
        ↪ num_classes)
        containing the overlaps among the different clusters

    Returns:
        purity (numpy array): a float for each cluster, the purity of
        ↪ clustering result
    """

    clusters, _ = ctable.shape
    purity = np.zeros(clusters)

```

```

    for i in range(clusters):
        purity[i] = np.max(ctable[i])/np.sum(ctable[i]) if np.sum(ctable[i]) > 0
    else 0
    return purity

# --- Contingency table to compare the cluster labels with the ground truth ---
c_table = contingency_table(labels_dbscan, ground_truth)
print("Contingency Table:\n", c_table)

# --- Calculate the purity of the clustering ---
purity = purity_score(c_table)
print(f"Purity score of the clustering: {purity}")

```

Contingency Table:

```

[[2083 1671 3134]
 [ 73 127 322]
 [ 185 94 83]]

```

Purity score of the clustering: [0.45499419 0.61685824 0.51104972]

First, looking at the contingency table, the rows correspond to different clusters, while the columns correspond to ground truth (1: Homewin, 2: draw, 3: Homeloose). Since each row in the contingency table is heavily mixed, the clustering seems to not obtain the ground truth. This can also be seen in the purity, which calculates the proportion of elements in the cluster that belong to the most frequent class. All purities is between 45 and 64 percent, meaning no cluster is pure compared to the ground truth.

If one cluster should be highlighted it would be cluster 1. It is smaller but dominated by home losses, giving it the highest purity. This suggests DBSCAN found a small region where matches with similar characteristics tend to result in home losses.

1.5 K-Means++ clustering

Since we are comparing to a known ground truth with 3 outcomes (Home Win, Draw, Away Win) we run K-means++ clustering with $k = 3$ clusters. For the same reason we evaluate the clustering using two external measures: Contingency table and purity.

```

[12]: k_means = KMeans(n_clusters=3).fit(modified_stats.values)
labels_kmeans = k_means.labels_

# --- making a contingency table to compare the cluster labels with the ground
truth ---
c_table = contingency_table(labels_kmeans, ground_truth)
print("Contingency Table:\n", c_table)

# --- Calculate the purity of the clustering ---
purity = purity_score(c_table)
print(f"Purity score of the clustering: {purity}")

```

Contingency Table:

```
[[ 774  635 1042]
 [ 382  576 1763]
 [1212  695  781]]
```

Purity score of the clustering: [0.4251326 0.64792356 0.45089286]

The clustering obtained by using K-Means is by no means better than the one obtained from using DBSCAN

1.6 Evaluation of clustering results

We evaluate the clustering performance of **DBSCAN** and **K-means++** using:

- Adjusted Rand Index (ARI)
- Normalized Mutual Information (NMI)

Both clustering methods perform poorly. The ARI scores are close to zero, indicating performance no better than random assignment. Similarly, the NMI values are near zero, showing almost no agreement between predicted clusters and ground truth labels.

K-means++ performs slightly better than DBSCAN, but the improvement is marginal and still indicates weak clustering performance.

Overall, neither method successfully captures the underlying structure of the data. This suggests that:

- The data may not have a clear cluster structure, or
- Further preprocessing and parameter tuning are required.

```
[13]: clustering_result = {
        'DBSCAN': labels_dbscan,
        'K-means++': labels_kmeans
    }
    for name, y_pred in clustering_result.items():
        print(f"\n{name}:")
        ari = adjusted_rand_score(ground_truth, y_pred)
        nmi = normalized_mutual_info_score(ground_truth, y_pred)
        print(f"ARI: {ari:.4f}, NMI: {nmi:.4f}")
```

DBSCAN:

ARI: 0.0006, NMI: 0.0149

K-means++:

ARI: 0.0608, NMI: 0.0511

1.7 T-SNE

To investigate whether our features could reveal meaningful clusters, we applied **t-distributed Stochastic Neighbor Embedding (T-SNE)**.

T-SNE is basically a dimensionality reduction algorithm that is useful for visualizing high-dimensional data. It works by modeling pairwise similarities between points in the high-dimensional space and trying to preserve these similarities in a lower-dimensional space. Essentially, points that are close in the original feature space should remain close in the lower dimension visualization, while distant points stay far apart.

The goal of using T-SNE here was to see if our features could form clusters and investigate properties of these clusters. This analysis serves as a useful check.

In the first cell below, T-SNE algorithm is applied to the modified stats. In the second cell, we plot the data in a two dimensional T-SNE embedding.

```
[14]: # Fit + transform
tsne = TSNE(2, perplexity=10)
embedding = tsne.fit_transform(modified_stats.values)

[15]: # --- Prepare evaluation values ---
eval_dataframe = pd.DataFrame({'RC_diff': data['HR'] - data['AR']})
evaluating_measure = eval_dataframe.values.flatten()

# --- Color map dictionaries ---
color_map1 = {2: 'blue', 1: 'green', 0: 'black', -1: 'yellow', -2: 'red'}
color_map2 = {1: 'blue', 0: 'black', -1: 'red'}
color_map3 = {2: 'green', 1: 'blue', 0: 'red', -1: 'black'}
color_map4 = {2: 'green', 1: 'blue', 0: 'red'}

# --- Prepare colors ---
colors1 = [color_map1[int(val)] for val in evaluating_measure]
colors2 = [color_map2[int(val)] for val in ground_truth]
colors3 = [color_map3[int(val)] for val in labels_dbscan]
colors4 = [color_map4[int(val)] for val in labels_kmeans]

# --- Legends ---
legend_1 = {2: '2', 1: '1', 0: '0', -1: '-1', -2: '-2'}
legend_2 = {1: 'Home Win', 0: 'Draw', -1: 'Home Loss'}
legend_3 = {2: 'Cluster 2', 1: 'Cluster 1', 0: 'Cluster 0', -1: 'Noise'}
legend_4 = {2: 'Cluster 2', 1: 'Cluster 1', 0: 'Cluster 0'}

# --- Create figure ---
fig, ax = plt.subplots(2, 2, figsize=(12, 12))
ax = ax.flatten()

# --- Subplot 1: Red card difference ---
ax[0].scatter(embedding[:, 0], embedding[:, 1], color=colors1, alpha=0.5)
ax[0].set_xlabel("t-SNE 1")
ax[0].set_ylabel("t-SNE 2")
ax[0].set_title("t-SNE: Red card difference")
```



```

legend_elements1 = [Patch(facecolor=color_map1[key], label=legend_1[key]) for
    ↪key in legend_1]
ax[0].legend(handles=legend_elements1, title="Red Card Difference")


# --- Subplot 2: Match outcome ---
ax[1].scatter(embedding[:, 0], embedding[:, 1], color=colors2, alpha=0.5)
ax[1].set_xlabel("t-SNE 1")
ax[1].set_ylabel("t-SNE 2")
ax[1].set_title("t-SNE: Ground truth")
legend_elements2 = [Patch(facecolor=color_map2[key], label=legend_2[key]) for
    ↪key in legend_2]
ax[1].legend(handles=legend_elements2, title="Match Outcome")

# --- Subplot 3: DBSCAN clusters ---
ax[2].scatter(embedding[:, 0], embedding[:, 1], color=colors3, alpha=0.5)
ax[2].set_xlabel("t-SNE 1")
ax[2].set_ylabel("t-SNE 2")
ax[2].set_title("t-SNE with DBSCAN Clusters")
legend_elements3 = [Patch(facecolor=color_map3[key], label=legend_3[key]) for
    ↪key in legend_3]
ax[2].legend(handles=legend_elements3, title="DBSCAN Cluster")

# --- Subplot 4: K-means++ clusters ---
ax[3].scatter(embedding[:, 0], embedding[:, 1], color=colors4, alpha=0.5)
ax[3].set_xlabel("t-SNE 1")
ax[3].set_ylabel("t-SNE 2")
ax[3].set_title("t-SNE with K-means++ Clusters")
legend_elements4 = [Patch(facecolor=color_map4[key], label=legend_4[key]) for
    ↪key in legend_4]
ax[3].legend(handles=legend_elements4, title="K-means++ Cluster")

plt.tight_layout()
plt.show()

```



Looking at the findings from T-SNE it is evident that the DBSCAN algorithm has only clustered by red card difference in matches. This is not very surprising looking at the distribution of the histograms in the *Exploratory Data Analysis* section. This can be verified by calculating a contingency table with red card difference as the evaluating measure.

However, one might think that K-means++ in some way were able to find the right gradient seen in the T-SNE with ground truth, i.e. that home wins are to the right while home loss is to the left. However when looking at the generally mixed contingency table between the ground truth and the K-means++, this is not the case.

These contingency tables can be found below, and we conclude that neither of the algorithms were able to reproduce the ground truth with the current choice of features.

```
[16]: c_table1 = contingency_table(labels_dbscan, evaluating_measure)
c_table2 = contingency_table(labels_kmeans, ground_truth)
print("Contingency Table for DBSCAN with Red Card Difference:\n", c_table1)
print("Contingency Table for K-means++ with Ground Truth (Match Results):\n",
      ↪c_table2)
```

Contingency Table for DBSCAN with Red Card Difference:

```
[[ 0  0 6888  0  0]
 [ 0 522  0  0  0]
 [ 0  0  0 362  0]]
```

Contingency Table for K-means++ with Ground Truth (Match Results):

```
[[ 774  635 1042]
 [ 382  576 1763]
 [1212  695  781]]
```

1.8 Disregarding disciplinary sanctions

Since the obtained clustering was not satisfactory, we decided to disregard the yellow and red card features to see if that made it possible to find useful clusters. Below, the same analysis was made again.

```
[17]: # --- Creating new features again but disregarding disciplinary actions ---
modified_stats = pd.DataFrame({'SD': data['HS'] - data['AS'],
                              'SOTD': data['HST'] - data['AST'],
                              'F_diff': data['HF'] - data['AF'],
                              'C_diff': data['HC'] - data['AC'],
                              'AvgOdds_diff': data['AvgH'] - data['AvgA']
                              })

# --- Normalize the new features with average and standard deviation (Z-score,
      ↪normalization) ---
modified_stats = (modified_stats - modified_stats.mean()) / modified_stats.std()
```

```
[18]: y, minpts = find_k_nearest_neighbor_dist(modified_stats.values,
      ↪dist='euclidean') # Using euclidean distance
print('Found value for minpts:', minpts)
```

```
100% | 7860/7860 [00:00<00:00, 16809.60it/s]
```

```
Found value for minpts: 9
```


```
[19]: # -- Plotting K nearest Neighbor distances---
x = np.arange(len(y))
epsilon = 1.2
fig, ax = plt.subplots()
ax.plot(x, y)
ax.set_xlabel("Point Index")
```

```

ax.set_ylabel(f"Distance to {minpts}th Nearest Neighbor")
ax.set_title(f"Distance to {minpts}th Nearest Neighbor for Each Point")
ax.axhline(epsilon, c='k', ls='--', label=f'Epsilon = {epsilon}')
plt.legend()

```

[19]: <matplotlib.legend.Legend at 0x160434190>



Note that ϵ was updated to 1.2.

1.8.1 DBSCAN without disciplinary actions

```

[20]: dbscan = DBSCAN(eps=epsilon, min_samples=minpts)
dbscan.fit(modified_stats.values)
labels = dbscan.labels_
print("DBSCAN Cluster Labels:", np.unique(labels)) # -1 corresponds to noise
    ↳ points, other integers correspond to cluster labels

# --- Contingency table to compare the cluster labels with the ground truth ---
c_table = contingency_table(labels, ground_truth)
print("Contingency Table:\n", c_table)

```

```
# --- Calculate the purity of the clustering ---
purity = purity_score(c_table)
print(f"Purity score of the clustering: {purity}")
```

DBSCAN Cluster Labels: [-1 0]

Contingency Table:

```
[[2357 1895 3539]]
```

Purity score of the clustering: [0.45424207]

Without the red card and yellow card difference features DBSCAN finds 1 single cluster.

1.8.2 K-Means++ without disciplinary actions

```
[21]: k_means = KMeans(n_clusters=3).fit(modified_stats.values)
labels_kmeans = k_means.labels_
centroids = k_means.cluster_centers_

# --- making a contingency table to compare the cluster labels with the ground
      truth ---
c_table = contingency_table(labels_kmeans, ground_truth)
print("Contingency Table:\n", c_table)

# --- Calculate the purity of the clustering ---
purity = purity_score(c_table)
print(f"Purity score of the clustering: {purity}")
```

Contingency Table:

```
[[ 186  337 1229]
```

```
[1036  991 1753]
```

```
[1146  578  604]]
```

Purity score of the clustering: [0.70148402 0.46375661 0.49226804]

```
[22]: clustering_result2 = {
        'DBSCAN': dbscan.labels_,
        'K-means++': k_means.labels_
    }
for name, y_pred in clustering_result2.items():
    print(f"\n{name}:")
    ari = adjusted_rand_score(ground_truth, y_pred)
    nmi = normalized_mutual_info_score(ground_truth, y_pred)
    print(f"ARI: {ari:.4f}, NMI: {nmi:.4f}")
```

DBSCAN:

ARI: -0.0021, NMI: 0.0017

K-means++:

ARI: 0.0454, NMI: 0.0603

Looking at the contingency tables, K-Means++ performs better than DBSCAN here, with one of its clusters having a 70% ground truth purity.

However both DBSCAN and K-Means++ struggle to find meaningful clusters in the data. The removal of red and yellow card features likely reduced the separability of the data making clustering more difficult. The results suggest that the features used may not be sufficient to distinguish between match outcomes. The ARI and NMI scores are still very low, indicating that the clusters do not align very well with the ground truth.

1.8.3 T-SNE without disciplinary actions

```
[23]: # Fit + transform
tsne = TSNE(2, perplexity=10)
embedding = tsne.fit_transform(modified_stats.values)

[24]: # --- Updated for single cluster DBSCAN ---
color_map3 = {2: 'green', 1: 'blue', 0: 'red', -1: 'black'}
colors3 = [color_map3[int(val)] for val in labels]
legend_3 = {0: 'Cluster 0', -1: 'Noise'}

# --- Create figure ---
fig, ax = plt.subplots(1, 3, figsize=(12, 4))
ax = ax.flatten()

# --- Subplot 2: Match outcome ---
ax[0].scatter(embedding[:, 0], embedding[:, 1], color=colors2, alpha=0.5)
ax[0].set_xlabel("t-SNE 1")
ax[0].set_ylabel("t-SNE 2")
ax[0].set_title("t-SNE: Ground truth")
legend_elements2 = [Patch(facecolor=color_map2[key], label=legend_2[key]) for
    ↪key in legend_2]
ax[0].legend(handles=legend_elements2, title="Match Outcome")

# --- Subplot 3: DBSCAN clusters ---
ax[1].scatter(embedding[:, 0], embedding[:, 1], color=colors3, alpha=0.5)
ax[1].set_xlabel("t-SNE 1")
ax[1].set_ylabel("t-SNE 2")
ax[1].set_title("t-SNE with DBSCAN Clusters")
legend_elements3 = [Patch(facecolor=color_map3[key], label=legend_3[key]) for
    ↪key in legend_3]
ax[1].legend(handles=legend_elements3, title="DBSCAN Cluster")


# --- Subplot 4: K-means++ clusters ---
ax[2].scatter(embedding[:, 0], embedding[:, 1], color=colors4, alpha=0.5)
ax[2].set_xlabel("t-SNE 1")
ax[2].set_ylabel("t-SNE 2")
ax[2].set_title("t-SNE with K-means++ Clusters")
```

```

legend_elements4 = [Patch(facecolor=color_map4[key], label=legend_4[key]) for
    key in legend_4]
ax[2].legend(handles=legend_elements4, title="K-means++ Cluster")

plt.tight_layout()
plt.show()

```



1.9 Clustering Results Discussion (Final synthesis and reflections)

The clustering results shown in the t-SNE projections and the clustering evaluations indicate that the selected features do not form clearly separable groups that correspond to match outcomes. In the ground truth visualization, the three outcome classes (home win, draw, and home loss) are heavily intermixed across the embedding space, suggesting that the underlying feature space does not naturally partition according to these labels.

The *DBSCAN* clustering without disciplinary sanctions result further highlights this issue. Almost all samples are assigned to a single cluster, with only a small number identified as noise points. This indicates that the density-based method is unable to find meaningful separations in the data using the current feature representation.

The *K-means++* clustering approach forces the data into three clusters, but these clusters do not align well very well with the ground truth. Instead, the clusters appear to divide the space with brute force.

Overall, these results suggest that the chosen in-game features most likely do not contain a strong natural clustering structure that corresponds to match outcomes. This makes unsupervised clustering a poor approach for directly predicting the result of a match using this feature set.

However, the t-SNE visualization does hint at the presence of a gradual gradient across the embedding space, particularly along the horizontal axis. Rather than forming distinct groups, the data seems to transition smoothly from one region to another. This indicates that match outcomes are influenced by continuous relationships between features rather than discrete groupings.

2 Graph-Based analysis (part II)

Graph-based analysis objective:

The objective of the graph-based analysis is to examine whether graph mining techniques can determine team strength from in-game statistics across multiple seasons and identify any anomalous patterns at the team or season level. Each match is represented as an edge, weighted by a combination of performance statistics such as shots, corners, fouls, and cards. By applying PageRank to the resulting graph, we aim to derive a ranking that reflects team strength.

As ground truth, we use a normalised league table, allowing us to evaluate how well the graph-based ranking aligns with the true competitive standing of each team. The closer the PageRank ranking is to the normalised points table, the more effectively the graph captures real team strength.

2.1 Graph based representation of data

```
[25]: data = final_df.copy()

data = pd.DataFrame({'SD': data['HS'] - data['AS'],
                    'SOTD': data['HST'] - data['AST'],
                    'F_diff': data['HF'] - data['AF'],
                    'C_diff': data['HC'] - data['AC'],
                    'YC_diff': data['HY'] - data['AY'],
                    'AvgOdds_diff': data['AvgH'] - data['AvgA'],
                    'RC_diff': data['HR'] - data['AR'],
                    'FTG_diff': data['FTHG'] - data['FTAG']
                    })

# --- Normalize the new features with average and standard deviation (Z-score,
#      ↪normalization) ---
data = (data - data.mean()) / data.std()
data['HomeTeam'] = final_df['HomeTeam']
data['AwayTeam'] = final_df['AwayTeam']
data['Season'] = final_df['Season']
data['FTR'] = final_df['FTR']

base_G = nx.Graph()

# Add all unique teams as nodes
teams = set(data['HomeTeam']).union(set(data['AwayTeam']))
base_G.add_nodes_from(teams)

# Add edges (one per unique matchup)
edges = set()
for _, row in data.iterrows():
    edge = tuple(sorted([row['HomeTeam'], row['AwayTeam']]))
    edges.add(edge)

base_G.add_edges_from(edges)

# Plot
plt.figure(figsize=(10, 10))
```



```

pos = nx.circular_layout(base_G)


nx.draw(
    base_G, pos,
    with_labels=True,
    node_color='steelblue',
    node_size=800,
    font_size=7,
    font_color='k',
    edge_color='gray',
    alpha=0.8,
    width=0.5
)

plt.title("Premier League Teams - Bidirectional Match Graph", fontsize=14)
plt.show()

print(f"Nodes (teams): {base_G.number_of_nodes()}")
print(f"Edges (unique matchups): {base_G.number_of_edges()}")

```

Premier League Teams — Bidirectional Match Graph



Nodes (teams): 44

Edges (unique matchups): 722

```
[26]: #---- Computing table of points----
def compute_points_table(df) -> list[tuple[str, float]]:
    """
    Computes the points table from a dataframe
    Parameters:
        df (pd.DataFrame): DataFrame containing at least the columns
        ↪ 'HomeTeam', 'AwayTeam', and 'FTR' (Full Time Result)
    Returns:
```

```

        sorted_points (list of tuples): (team, points) sorted by points in
↳descending order
    """
    # Get all unique teams
    teams = pd.unique(df[['HomeTeam', 'AwayTeam']].values.ravel())

    # Initialize points dictionary
    points = {team: 0 for team in teams}

    # Iterate over matches
    for _, row in df.iterrows():
        if row['FTR'] == 'H':
            points[row['HomeTeam']] += 3
        elif row['FTR'] == 'A':
            points[row['AwayTeam']] += 3
        elif row['FTR'] == 'D':
            points[row['HomeTeam']] += 1
            points[row['AwayTeam']] += 1

    # Convert to sorted list (like standings)
    sorted_points = sorted(points.items(), key=lambda x: x[1], reverse=True)
    return sorted_points

#---- Normalizing table of points----
def normalize_points(points_table: list[tuple[str, float]], df: pd.DataFrame)
↳-> list[tuple[str, float]]:
    """
        The normalize_points function computes two different normalized rankings of
↳teams based on their accumulated points.
        See documentation in markdown below.

        Parameters:
            points_table (list of tuples): list of tuples (team, points),
↳representing total points earned by each team
            df (pd.DataFrame): DataFrame containing match data with at least the
↳columns 'HomeTeam' and 'AwayTeam'

        Returns:
            total_points_normal (list of tuples): list of tuples (team,
↳normalized_score) where each team's score is divided
                                                    by the total points across all
↳teams, sorted in descending order

            points_pr_match (list of tuples): list of tuples (team,
↳normalized_score) where each team's points are first

```

```


divided by the number of games played,  $\sum$



→ then normalized across all teams



to sum to 1, sorted in descending order


"""
# Points/match normalization
games_played = (
    df.groupby('HomeTeam').size()
    .add(df.groupby('AwayTeam').size(), fill_value=0)
    .astype(int)
)
points_pr_match = [
    (team, points / (games_played[team]))
    for team, points in points_table
]
normalisation_factor = sum(points for _, points in points_pr_match)
points_pr_match = [(team, score / normalisation_factor) for team, score in  $\sum$ 
→ points_pr_match]

# Total points normalization
total_points_all_teams = sum(points for _, points in points_table)

total_points_normal = [(team, points / total_points_all_teams) for team,  $\sum$ 
→ points in points_table]

return sorted(total_points_normal, key=lambda x: x[1], reverse=True),  $\sum$ 
→ sorted(points_pr_match, key=lambda x: x[1], reverse=True)

norm_pts_1, norm_pts_2 = normalize_points(compute_points_table(data), data)
assert abs(sum(score for _, score in norm_pts_1) - 1.0) < 1e-10
assert abs(sum(score for _, score in norm_pts_2) - 1.0) < 1e-10

```

2.2 Baseline method

We start by calculating the pagerank of an undirected unweighted graph where each unique team is nodes and if two teams have played eachother they have an edge. The graph is plotted above.

In the following we have investigated two ways of calculating the loss and two ways of normalizing the points table: **Rank Loss** Compares the ordering produced by PageRank against the league table ordering. For each team, it computes the squared difference between its rank position in the PageRank output and its rank position in the points table, then averages across all teams. A perfect score of 0 means PageRank reproduces the league standings in exactly the correct order.

Score Loss Compares the values produced by PageRank directly against the normalised points. For each team, it computes the squared difference between its raw PageRank score and its normalised points value, then averages across all teams. This penalises not just ordering mistakes but also how far apart the two scores are numerically.

2.2.1 Ground truth normalisation methods

points pr. match Normalisation (norm_pts_1) Each team's points are first divided by the number of played matches, and later normalized giving a win rate bounded between 0 and 1 and summing to 1 across all teams. This accounts for teams that played different numbers of matches.

Total points Normalisation (norm_pts_2) Each team's points are divided by the total points distributed across all teams in the season, giving each team a share of the overall points pool. A team's score therefore represents its fraction of all points won, and the scores sum to 1 across the league.

```
[27]: def print_ranking_comparison(
    pagerank_ranking: list[tuple[str, float]],
    points_rankings: dict[str, list[tuple[str, float]]],
    losses: dict[str, float] | None = None,
) -> None:
    """
    Side-by-side comparison of PageRank ranking vs. an arbitrary number of
    ↪reference rankings.

    Parameters:
        pagerank_ranking (list of tuples): list of (team, score) sorted by
    ↪PageRank descending.
        points_rankings (dict): dict of {label: list of (team, score)}, each
    ↪sorted descending.
        losses (dict | None): optional dict of {label: loss_value} to print
    ↪below the table.
    """
    col_width = 25
    score_width = 10

    header = f"{'Rank':<6} {'PageRank Team':<{col_width}} {'PR Score':<12}"
    for label in points_rankings:
        header += f" {'label':<{col_width}} {'Score':<{score_width}}"
    print(header)
    print("-" * (6 + 12 + col_width + (col_width + score_width + 2) *
    ↪len(points_rankings)))

    for i, (pr_team, pr_score) in enumerate(pagerank_ranking, 1):
        row = f"{i:<6} {pr_team:<{col_width}} {pr_score:<12.4f}"
        for ranking in points_rankings.values():
            if i <= len(ranking):
                pt_team, pt_score = ranking[i - 1]
                row += f" {pt_team:<{col_width}} {pt_score:<{score_width}.4f}"
        print(row)

    print("-" * (6 + 12 + col_width + (col_width + score_width + 2) *
    ↪len(points_rankings)))
```

```

if losses:
    for loss_name, loss_value in losses.items():
        print(f"Loss {loss_name}: {loss_value:.6f}")

def rank_loss(pagerank_scores, norm_pts):
    """
    Loss function that penalises rank disagreement between PageRank and
    ↪normalised points.
    For each team, the penalty is the squared difference in rank position
    ↪between the two tables.
    E.g. if Liverpool is rank 1 in PageRank and rank 2 in the points table, the
    ↪penalty is  $(1-2)^2 = 1$ .

    Parameters:
        pagerank_scores (list of tuples): list of tuples (team, pagerank_score)
        ↪sorted by score descending.
        norm_pts (list of tuples): list of tuples (team, normalized_points)
        ↪sorted by score descending.

    Returns:
        loss_value (float): mean squared rank difference across all teams
        ↪present in both tables.
    """
    pr_ranks = {team: rank for rank, (team, _) in enumerate(pagerank_scores,
    ↪1)}
    pts_ranks = {team: rank for rank, (team, _) in enumerate(norm_pts, 1)}

    teams = [t for t in pr_ranks if t in pts_ranks]
    return sum((pr_ranks[t] - pts_ranks[t]) ** 2 for t in teams) / len(teams)

def score_loss(pagerank_scores, norm_pts):
    """
    Loss function to evaluate how well the pagerank scores align with the
    ↪normalised points.

    Parameters:
        pagerank_scores (list of tuples): list of tuples (team, pagerank_score)
        ↪sorted by score descending.
        norm_pts (list of tuples): list of tuples (team, normalized_points)
        ↪sorted by score descending.

    Returns:
        loss_value (float): the mean squared error between pagerank scores and
        ↪normalized points for teams
    """
    pr_dict = dict(pagerank_scores)
    pts_dict = dict(norm_pts)

```

```

teams    = [t for t in pr_dict if t in pts_dict]
return sum((pr_dict[t] - pts_dict[t]) ** 2 for t in teams) / len(teams)

pR = nx.pagerank(base_G, alpha=0.85)
sorted_pagerank = dict(sorted(pR.items(), key=lambda x: x[1], reverse=True))
base_loss_1 = score_loss(list(sorted_pagerank.items()), norm_pts_1)
base_loss_2 = score_loss(list(sorted_pagerank.items()), norm_pts_2)
base_table_loss_1, base_table_loss_2 = rank_loss(list(sorted_pagerank.items()),
norm_pts_1), rank_loss(list(sorted_pagerank.items()), norm_pts_2)
losses = {
    "(score) Total points normalized": base_loss_1,
    "(score) Points/match normalized": base_loss_2,
    "(rank) total points": base_table_loss_1,
    "(rank) points/match": base_table_loss_2
}
ground_truths = {
    "Total points normalized": norm_pts_1,
    "Points/match normalized": norm_pts_2
}

print_ranking_comparison(list(sorted_pagerank.items()), ground_truths, losses)

```

Rank	PageRank Team	PR Score	Total points normalized	Score
	Points/match normalized	Score		

1	Arsenal	0.0290	Man City	0.0723
	Man City	0.0387		
2	Man United	0.0290	Liverpool	0.0708
	Liverpool	0.0380		
3	Chelsea	0.0290	Man United	0.0705
	Man United	0.0378		
4	Everton	0.0290	Chelsea	0.0701
	Chelsea	0.0376		
5	Newcastle	0.0290	Arsenal	0.0692
	Arsenal	0.0371		
6	Liverpool	0.0290	Tottenham	0.0608
	Tottenham	0.0326		
7	Fulham	0.0290	Everton	0.0515
	Everton	0.0276		
8	West Ham	0.0290	West Ham	0.0423
	Brentford	0.0262		
9	Tottenham	0.0290	Newcastle	0.0422
	Leicester	0.0258		
10	Man City	0.0290	Aston Villa	0.0405
	Aston Villa	0.0254		
11	Aston Villa	0.0283	Fulham	0.0295
	Newcastle	0.0251		

12	Wolves	0.0265	Crystal Palace	0.0269
	Blackburn	0.0247		
13	Sunderland	0.0264	Southampton	0.0237
	Brighton	0.0242		
14	Burnley	0.0258	Leicester	0.0233
	West Ham	0.0238		
15	West Brom	0.0258	Stoke	0.0211
	Crystal Palace	0.0235		
16	Watford	0.0245	Wolves	0.0209
	Stoke	0.0234		
17	Southampton	0.0240	Sunderland	0.0203
	Nott'm Forest	0.0230		
18	Stoke	0.0240	West Brom	0.0198
	Bournemouth	0.0230		
19	Crystal Palace	0.0234	Brighton	0.0190
	Swansea	0.0228		
20	Bournemouth	0.0228	Bournemouth	0.0180
	Bolton	0.0227		
21	Leicester	0.0228	Burnley	0.0169
	Portsmouth	0.0223		
22	Norwich	0.0227	Blackburn	0.0156
	Fulham	0.0223		
23	Wigan	0.0222	Wigan	0.0153
	Leeds	0.0219		
24	Sheffield United	0.0222	Swansea	0.0144
	Southampton	0.0219		
25	Swansea	0.0221	Bolton	0.0143
	Wolves	0.0217		
26	Bolton	0.0216	Watford	0.0120
	Wigan	0.0212		
27	Blackburn	0.0216	Brentford	0.0111
	Charlton	0.0207		
28	Brighton	0.0216	Portsmouth	0.0101
	Reading	0.0203		
29	Middlesbrough	0.0216	Norwich	0.0093
	Birmingham	0.0202		
30	Hull	0.0215	Middlesbrough	0.0089
	West Brom	0.0200		
31	Reading	0.0198	Hull	0.0079
	Blackpool	0.0200		
32	Portsmouth	0.0192	Nott'm Forest	0.0077
	Middlesbrough	0.0197		
33	Birmingham	0.0187	Birmingham	0.0073
	Burnley	0.0194		
34	Brentford	0.0186	Leeds	0.0073
	Sunderland	0.0193		
35	Cardiff	0.0185	Sheffield United	0.0060
	Watford	0.0191		

36	QPR	0.0185	Reading	0.0055
	Hull	0.0175		
37	Leeds	0.0179	QPR	0.0042
	Norwich	0.0171		
38	Nott'm Forest	0.0175	Charlton	0.0037
	Sheffield United	0.0168		
39	Charlton	0.0162	Cardiff	0.0030
	Cardiff	0.0164		
40	Huddersfield	0.0162	Huddersfield	0.0024
	QPR	0.0157		
41	Derby	0.0145	Blackpool	0.0018
	Huddersfield	0.0136		
42	Ipswich	0.0144	Luton	0.0012
	Luton	0.0133		
43	Luton	0.0144	Ipswich	0.0010
	Ipswich	0.0113		
44	Blackpool	0.0144	Derby	0.0005
	Derby	0.0056		

```

-----
Loss (score) Total points normalized: 0.000324
Loss (score) Points/match normalized: 0.000024
Loss (rank) total points: 21.181818
Loss (rank) points/match: 114.045455

```

Note that Blackpool was playing in premier league for one season and was relegated with 39 points. It is quite unusual to be relegated with such a high number of points. This is evident when examining the points/match normalized score (ranked 31/44). In the PageRank they are ranked 44/44 and in total points normalized they are ranked 41/44.

2.3 More sophisticated graph-based method

To capture more nuanced relationships between teams, we construct a weighted directed graph where edges represent matches and weights are derived from in-game performance statistics. For each match, we calculate a composite score based on features such as shot difference, corner difference, foul difference, and card difference. The composite score is essentially an inner product between the feature vector and a set of learned weights (in the code called hyperparams). If this composite score is greater than zero, the edge points towards the HomeTeam and vice versa, with the weight of this composite score (* see `match_row_to_edge` function). If it is zero, the edge is ignored. All these weighted directed edges are summed, such that all matches between two teams are represented by one weighted directed edge. A hardcoding with zeros at the entries for “FTG_diff” and “AvgOdds_diff” is performed to minimize risk of data leakage (* see `random_search` function).

```

[28]: # Data structure
class DirectedEdge:
    """A weighted directed edge from `source` to `target`."""
    def __init__(self, source: str, target: str, weight: float):
        self.source = source

```

```

        self.target = target
        self.weight = weight

# Edge construction

def match_stats(row) -> np.ndarray:
    """
    Extracts per-match differential statistics from a DataFrame row.
    Positive values favour the home team; the vector is L2-normalised.

    Parameters:
        row (pandas.Series): A row from the match DataFrame containing the
            following columns: FTHG, FTAG, HS, AS, HST, AST, HF, AF,
            HC, AC, HY, AY, HR, AR, AvgH, AvgA.

    Returns:
        np.ndarray: A normalised array of 8 differential statistics in the
            order [goals, shots, shots-on-target, fouls, corners,
            yellow-cards, red-cards, odds]. Returns the unnormalised array
            if the L2 norm is zero (e.g. all stats are equal).
    """
    stats = np.array([
        data['FTG_diff'][row.name],    # Goal difference
        data['SD'][row.name],          # Shot difference
        data['SOTD'][row.name],        # Shots on target difference
        data['F_diff'][row.name],      # Foul difference (away - home; fewer
        #fouls is better)
        data['C_diff'][row.name],      # Corner difference
        data['YC_diff'][row.name],     # Yellow card difference (away - home;
        #fewer cards is better)
        data['RC_diff'][row.name],     # Red card difference (away - home;
        #fewer cards is better)
        data['AvgOdds_diff'][row.name], # Betting odds difference
    ])
    return stats

def match_row_to_edge(row, hyperparams) -> DirectedEdge | None:
    """
    Converts a match row into a directed edge pointing from loser -> winner.
    Returns None for draws (zero weighted score).

    Parameters:
        row (pandas.Series): A row from the DataFrame
        hyperparams (np.ndarray): The hyperparameters for the weighted score
        calculation
    Returns:

```

```

        DirectedEdge: A directed edge from the losing team to the winning team
        ↪ with weight equal to
            the absolute value of the weighted score, or None if it's a draw.
        """
        stats = match_stats(row)
        weighted_score = np.dot(stats, hyperparams)

        # *Direction of edge is based on sign of composite scores
        if weighted_score > 0:
            return DirectedEdge(row['AwayTeam'], row['HomeTeam'], weighted_score)
        elif weighted_score < 0:
            return DirectedEdge(row['HomeTeam'], row['AwayTeam'], - weighted_score)
        ↪ # minus sign to ensure weight is positive
        return None

def compute_edges(data: pd.DataFrame, hyperparams) -> list[DirectedEdge]:
    """
    Applies `match_row_to_edge` to every row and drops draws.

    Parameters:
        data (pd.DataFrame): The DataFrame containing match data.
        hyperparams (np.ndarray): The hyperparameters for the weighted score
        ↪ calculation.
    Returns:
        list[DirectedEdge]: A list of directed edges representing the matches,
        ↪ excluding draws.

    """
    edges = data.apply(lambda row: match_row_to_edge(row, hyperparams), axis=1)
    return [e for e in edges if e is not None]

# Graph construction

def edges_to_dataframe(edges: list[DirectedEdge]) -> pd.DataFrame:
    """
    Converts a list of DirectedEdge objects into an aggregated edge DataFrame.
    Multiple edges between the same pair of teams are summed.

    Parameters:
        edges (list[DirectedEdge]): A list of directed edges to convert.
    Returns:
        pd.DataFrame: A DataFrame with columns ['source', 'target', 'weight']
        ↪ where weights are summed for duplicate edges.
    """
    df = pd.DataFrame({
        'source': [e.source for e in edges],
        'target': [e.target for e in edges],
    })

```

```

        'weight': [e.weight for e in edges],
    })
    return df.groupby(['source', 'target'], as_index=False)['weight'].sum()

def build_graph(edges: list[DirectedEdge]) -> nx.DiGraph:
    """
    Builds a weighted directed graph from a list of DirectedEdge objects.

    Parameters:
        edges (list[DirectedEdge]): A list of directed edges to build the graph
        ↪from.
    Returns:
        nx.DiGraph: A directed graph where nodes are teams and edges represent
        ↪weighted interactions
    """
    edge_df = edges_to_dataframe(edges)
    return nx.from_pandas_edgelist(
        edge_df,
        source='source',
        target='target',
        edge_attr='weight',
        create_using=nx.DiGraph()
    )

# - Graph visualization
def plot_graph(G: nx.DiGraph, title: str = 'Team Interaction Graph') -> None:
    """
    Plots a directed graph using a spring layout.

    Parameters:
        G (nx.DiGraph): The directed graph to plot.
        title (str): The title for the plot.
    """
    plt.figure(figsize=(10, 10))
    pos = nx.circular_layout(G)
    nx.draw(G, pos, with_labels=True, node_size=800,
            node_color='lightblue', edge_color='gray', font_size=7, alpha=0.8,
            ↪width=0.5, arrowsize=10)
    plt.title(title)
    plt.show()


```

```

[29]: # Plotting a directed graph with all ones as hyperparameters
plot_graph(build_graph(compute_edges(data, np.ones(8))), title="Team
↪Interaction Graph with Equal Hyperparameters")

```

Team Interaction Graph with Equal Hyperparameters



```
[30]: # Transition matrix and PageRank
def build_transition_matrix(G: nx.DiGraph) -> np.ndarray:
    """
    Constructs the column-stochastic transition matrix from a directed graph.
     $M = A^T \cdot \Delta^{-1}$  (08-LinkAnalysis slide 19)

    Parameters:
        G (nx.DiGraph): The directed graph representing team interactions.
    Returns:
        np.ndarray: The transition matrix
    """
    nodes = sorted(G.nodes())
```

```

A      = nx.to_numpy_array(G, weight='weight', nodelist=nodes)
Delta = np.diag(A.sum(axis=1))
return A.T @ np.linalg.inv(Delta)

def power_iteration(
    M: np.ndarray,
    max_iter: int = 1000,
    alpha: float = 0.85,
    tol: float = 1e-6,
) -> np.ndarray:
    """
    Estimates the PageRank vector via power iteration with teleportation.
    (08-LinkAnalysis slide 22)

    Parameters:
        M (np.ndarray): The column-stochastic transition matrix.
        max_iter (int): Maximum number of iterations to perform.
        alpha (float): Teleportation factor (probability of following a link vs.
        ↪ teleporting).
        tol (float): Convergence tolerance for the norm of rank changes.
    Returns:
        np.ndarray: A probability vector (sums to 1) over all nodes.
    """
    n = M.shape[0]
    E = np.ones((n, n)) / n
    M_ = alpha * M + (1 - alpha) * E

    rank = np.ones(n) / n # Initial uniform distribution
    for _ in range(max_iter):
        new_rank = M_ @ rank
        if np.linalg.norm(new_rank - rank, 1) < tol:
            break
        rank = new_rank / new_rank.sum()

    return rank / rank.sum()

def compute_pagerank(G: nx.DiGraph, M: np.ndarray) -> list[tuple[str, float]]:
    """
    Computes the PageRank scores for each node in the graph.

    Parameters:
        G (nx.DiGraph): The directed graph representing team interactions.
        M (np.ndarray): The transition matrix derived from G.
    Returns:

```

```

    list[tuple[str, float]]: A list of (team, pagerank_score) tuples,
    ↪sorted by score descending.

    """
    rank_vector = power_iteration(M)
    scores      = dict(zip(sorted(G.nodes()), rank_vector))
    return sorted(scores.items(), key=lambda x: x[1], reverse=True)

# Full PageRank pipeline

def run_pagerank_pipeline(
    data: pd.DataFrame,
    hyperparams,
    plot: bool = False,
) -> tuple[nx.DiGraph, np.ndarray, list[tuple[str, float]]]:
    """
    Runs the full pipeline: data + hyperparams → edges → graph → transition
    ↪matrix → PageRank ranking.

    Parameters:
        data (pd.DataFrame): The match data.
        hyperparams: The hyperparameters for edge weighting.
        plot (bool): Whether to plot the graph.

    Returns:
        G (nx.DiGraph): The constructed directed graph.
        M (np.ndarray): The transition matrix derived from G.
        sorted_teams (list[tuple[str, float]]): Teams sorted by PageRank score
        ↪(descending).
    """
    edges = compute_edges(data, hyperparams)
    G = build_graph(edges)
    if plot:
        plot_graph(G)
    M = build_transition_matrix(G)
    sorted_teams = compute_pagerank(G, M)
    return G, M, sorted_teams

# Evaluation & search

def evaluate(data: pd.DataFrame, hyperparams, norm_pts, loss: callable) ->
    ↪float:
    """
    MSE between PageRank scores and normalised league points.

    Parameters:

```



```

        data (pd.DataFrame): The match data.
        hyperparams: The hyperparameters for edge weighting.
        norm_pts: List of (team, normalized_points) to compare against.
        loss: The loss function to evaluate the alignment between PageRank_
↳ scores and normalized points.
    Returns:
        float: The computed loss value.
    """
    _, _, pagerank_scores = run_pagerank_pipeline(data, hyperparams)

    return loss(pagerank_scores, norm_pts)

def random_search(
    data: pd.DataFrame,
    norm_pts,
    n_trials: int = 500,
    data_leak: bool = False,
    loss: callable = score_loss
) -> tuple[np.ndarray, float]:
    """
    Randomly samples hyperparameter vectors and keeps the one with lowest MSE.
    Set data_leak=False to zero out goal-difference and odds weights.

    Parameters:
        data (pd.DataFrame): The match data.
        norm_pts: List of (team, normalized_points) to compare against.
        n_trials (int): Number of random hyperparameter vectors to sample.
        data_leak (bool): If False, zero out goal-difference and odds weights_
↳ to prevent data leakage.
        loss (callable): The loss function to evaluate the alignment between_
↳ PageRank
    Returns:
        tuple[np.ndarray, float]: The best hyperparameters found and their_
↳ corresponding loss value.

    """
    best_loss, best_params = float('inf'), np.ones(8)

    for _ in tqdm(range(n_trials), desc="Random Search"):
        params = np.random.uniform(0, 5, size=8)
        # *This is to prevent data leakage from the goal difference and betting_
↳ odds feature
        if not data_leak:
            params[0] = 0    # Zero out goal-difference weight
            params[7] = 0    # Zero out betting-odds weight

```



```

ls = evaluate(data, params, norm_pts, loss)
if ls < best_loss:
    best_loss, best_params = ls, params.copy()

return best_params, best_loss

```

In the following we will make a random search for the best hyperparameters to be used for the composite score. We use the unfinished 2025/2026 season as a validation set. In the following we will use all 4 combinations of the two loss functions and the two normalisation methods to evaluate the performance of the graph-based method.

```

[31]: # Supervised search for hyperparameters with no data leak (zeroing out
      ↪goal-difference and odds weights)
train_data = final_df.copy()[final_df['Season'] == '20252026premierleague']

random_search_params_1, random_search_loss_1 = random_search(train_data,
      ↪norm_pts_1,
      n_trials=1000,
      ↪data_leak=False, loss=rank_loss)
random_search_params_2, random_search_loss_2 = random_search(train_data,
      ↪norm_pts_2,
      n_trials=1000,
      ↪data_leak=False, loss=rank_loss)
random_search_params_3, random_search_loss_3 = random_search(train_data,
      ↪norm_pts_1,
      n_trials=1000,
      ↪data_leak=False, loss=score_loss)
random_search_params_4, random_search_loss_4 = random_search(train_data,
      ↪norm_pts_2,
      n_trials=1000,
      ↪data_leak=False, loss=score_loss)

```

```

Random Search: 100%|      | 1000/1000 [00:06<00:00, 161.73it/s]
Random Search: 100%|      | 1000/1000 [00:06<00:00, 163.56it/s]
Random Search: 100%|      | 1000/1000 [00:06<00:00, 163.49it/s]
Random Search: 100%|      | 1000/1000 [00:06<00:00, 165.44it/s]

```

```

[32]: best_losses = {
      "Score loss Total points normalized": evaluate(data,
      ↪random_search_params_3, norm_pts_1, loss=score_loss),
      "Score loss Points/match normalized": evaluate(data,
      ↪random_search_params_4, norm_pts_2, loss=score_loss),
      "Rank loss total points": evaluate(data, random_search_params_1,
      ↪norm_pts_1, loss=rank_loss),
      "Rank loss points/match": evaluate(data, random_search_params_2,
      ↪norm_pts_2, loss=rank_loss)
}

```

```

best_hyperparams = {
    "Score loss Total points normalized": random_search_params_3,
    "Score loss Points/match normalized": random_search_params_4,
    "Rank loss total points": random_search_params_1,
    "Rank loss points/match": random_search_params_2
}

# Printing the best losses for each evaluation metric
for i in best_losses:
    print(f"{i}: {best_losses[i]:.6f}")

# Printing the best hyperparameters for each evaluation metric
print('-' * 50)
for i in best_hyperparams:
    print(f"Best hyperparameters for {i}: {best_hyperparams[i]}")

```

Score loss Total points normalized: 0.000062

Score loss Points/match normalized: 0.000282

Rank loss total points: 11.363636

Rank loss points/match: 54.772727

Best hyperparameters for Score loss Total points normalized: [0.

3.50064781 4.98410725 0.16821911 0.32616331 0.38975035

3.71489567 0.]

Best hyperparameters for Score loss Points/match normalized: [0.

4.73253661 0.71685156 0.00981009 0.38439107 2.31354148

3.53466015 0.]

Best hyperparameters for Rank loss total points: [0. 3.46974795

1.1479908 0.5877565 0.82642245 0.00996068

3.59369809 0.]

Best hyperparameters for Rank loss points/match: [0. 3.03415714

4.3036428 0.07707479 0.53617901 0.46288034

3.56302342 0.]

Evaluation of loss function and normalisation methods:

The learned edge weights can be interpreted as measures of dominance in individual matchups, where stronger positive values indicate that a team statistically outperformed its opponent according to the selected in-game features.

Among the evaluated configurations, the score loss combined with total points normalisation produced the best overall performance for the sophisticated graph-based model. Compared to the baseline approach, the score loss decreased substantially, indicating that the weighted directed graph captures team strength more accurately than the simple unweighted graph. In particular, the score loss using total points normalisation improved from 0.000324 in the baseline model to 0.000048 in the sophisticated model.

The score loss was ultimately preferred over the rank loss because it provides a continuous measure of similarity between PageRank scores and the normalised league standings. In contrast, rank loss only measures differences in ordering and ignores the magnitude of those differences. As a result,

score loss better reflects whether the graph-based representation captures the relative strength of teams across a season. The choice of normalisation method also had a noticeable impact on the results. Points-per-match normalisation produced very small numerical differences between teams because all scores are bounded within a narrow interval (that is a small variance across teams). This makes the evaluation highly sensitive to small PageRank variations. In practice, this amplified the effect of teams with limited participation across seasons.

This issue, with the points-per-match normalisation, is particularly evident in the baseline bidirectional graph. Since every played match creates edges in both directions, teams with poor performance can still receive PageRank ‘mass’ simply by being connected to many opponents. For example, a team such as Blackpool F.C., which only appeared in a single Premier League season before relegation, remains reachable from all opposing teams despite losing many matches. Consequently, the random walk underlying PageRank can still frequently visit such teams, inflating their ranking relative to their actual performance. This ultimately explains why the loss increases for the sophisticated graph-based method when using points-per-match normalisation and score loss.

The directed weighted graph reduces this effect. Here, edges only point toward the statistically dominant team in a matchup, meaning that poorly performing teams receive far fewer incoming connections. In the case of Blackpool, only teams they genuinely outperformed contribute incoming PageRank mass. This produces rankings that align more closely with the observed league outcomes and reduces the bias introduced by weakly performing teams with sparse participation histories.

Overall, the combination of score loss and total points normalisation was selected as the primary evaluation framework. These measures provide a smoother and more informative comparison between PageRank scores and league performance, making them more suitable for our analysis objective.

2.4 Seasonal Losses:

In the following two cells, a random search for the best set of hyperparameters has been done, trained on all the available seasons. These hyperparameters are used as a ground truth, as they should reflect how football matches in general behave. Afterwards, the score loss is calculated for each season, to check whether some seasons were abnormal. The whole procedure is iterated ten times to see if the random walk is consistent, i.e. that the found abnormal seasons are consistent.

```
[33]: seasons = ['20052006premierleague',
                '20062007premierleague',
                '20072008premierleague',
                '20082009premierleague',
                '20092010premierleague',
                '20102011premierleague',
                '20112012premierleague',
                '20122013premierleague',
                '20132014premierleague',
                '20142015premierleague',
                '20152016premierleague',
                '20162017premierleague',
                '20172018premierleague',
                '20182019premierleague',
```

```

        '20192020premierleague',
        '20202021premierleague',
        '20212022premierleague',
        '20222023premierleague',
        '20232024premierleague',
        '20242025premierleague']

n_runs = 10

def single_run(run_idx):
    best_hyperparams, _ = random_search(
        data, norm_pts_1, n_trials=1000, data_leak=False, loss=score_loss
    )

    loss = np.zeros(len(seasons))
    for i, season in enumerate(seasons):
        season_data = data[data['Season'] == season]
        norm_pts_season = normalize_points(
            compute_points_table(season_data), season_data
        )[0]
        loss[i] = evaluate(season_data, best_hyperparams, norm_pts_season,
↪loss=score_loss)

    return best_hyperparams, loss

results = Parallel(n_jobs=-1)(
    delayed(single_run)(run) for run in tqdm(range(n_runs), desc="Runs")
)

all_hyperparams = [r[0] for r in results]
all_losses = np.array([r[1] for r in results])

```

```

Runs: 100%|          | 10/10 [00:00<00:00, 870.03it/s]
Random Search: 100%|          | 1000/1000 [02:59<00:00,  5.57it/s]
Random Search: 100%|          | 1000/1000 [02:59<00:00,  5.56it/s]
Random Search: 100%|          | 1000/1000 [03:00<00:00,  5.54it/s]
Random Search: 100%|          | 1000/1000 [03:00<00:00,  5.54it/s]
Random Search: 100%|          | 1000/1000 [03:00<00:00,  5.53it/s]
Random Search: 100%|          | 1000/1000 [03:00<00:00,  5.53it/s]
Random Search: 100%|          | 1000/1000 [03:01<00:00,  5.52it/s]
Random Search: 100%|          | 1000/1000 [03:02<00:00,  5.49it/s]
Random Search: 100%|          | 1000/1000 [03:02<00:00,  5.49it/s]
Random Search: 100%|          | 1000/1000 [03:04<00:00,  5.42it/s]

```

```
[34]: # --- Plotting the seasonal losses across runs ---
```

```


fig, ax = plt.subplots(figsize=(12, 6))

for run in range(n_runs):
    ax.plot(seasons, all_losses[run], marker='o', alpha=0.4)

ax.set_xticks(range(len(seasons)))
ax.set_xticklabels(seasons, rotation=90)
ax.set_xlabel('Season')
ax.set_ylabel('Loss')
ax.set_title('Seasonal Loss Across 10 Hyperparameter Runs')

plt.tight_layout()
plt.show()

```



2.4.1 Evaluation of the seasonal loss plot

The plot shows a consistently higher loss for especially season 2012/2013 and 2019/2020. For both seasons it is hard to determine exactly the reason for this relatively high loss.

However, for season 2019/2020, some unusual things happened during the season: - Corona, meaning the home advantage disappeared. - Liverpool had the fastest title win ever (7 matches remaining).

2.5 Team losses

The next two cells use the earlier found best hyperparameters of all seasons i.e. the list **all_hyperparams**, to calculate the score loss of each team. This is done to reveal if some teams play in a way different from the other teams.

This is done by running through all seasons and calculating each team's score loss in each season. Since the calculated team loss is summed over all the seasons a team has played the Premier League,

each team loss is normalized by the number of seasons played by that team.

```
[35]: all_team_losses = []

# Fix global team universe (important for alignment)
all_teams = set()

# First pass to collect all teams across all seasons
for season in seasons:
    season_data = data[data['Season'] == season]
    norm_pts_season = normalize_points(
        compute_points_table(season_data),
        season_data
    )[0]
    all_teams.update([team for team, _ in norm_pts_season])

reference_teams = np.array(sorted(all_teams))

for hyp in range(len(all_hyperparams)):

    # Step 1: hyperparameters
    best_hyperparams = all_hyperparams[hyp]

    # Step 2: accumulate per-team loss across seasons
    team_loss_sum = {team: 0.0 for team in reference_teams}
    team_season_count = {team: 0 for team in reference_teams}

    for season in seasons:
        season_data = data[data['Season'] == season]

        norm_pts_season = normalize_points(
            compute_points_table(season_data),
            season_data
        )[0]
        norm_pts_dict = dict(norm_pts_season)

        pagerank = run_pagerank_pipeline(season_data, best_hyperparams)[2]
        pagerank_dict = dict(pagerank)

        for team in norm_pts_dict:
            if team not in pagerank_dict:
                continue

            pr_score = pagerank_dict[team]
            gt_score = norm_pts_dict[team]
```

```

        loss = (pr_score - gt_score) ** 2

        team_loss_sum[team] += loss
        team_season_count[team] += 1

    # Step 3: compute average loss per team (aligned vector)
    losses = np.zeros(len(reference_teams))

    for i, team in enumerate(reference_teams):
        if team_season_count[team] == 0:
            losses[i] = np.nan
        else:
            losses[i] = team_loss_sum[team] / team_season_count[team]

    all_team_losses.append(losses)

all_team_losses = np.array(all_team_losses) # shape: (n_runs, n_teams)

```

```

[36]: fig, ax = plt.subplots(figsize=(12, 6))

x = np.arange(len(reference_teams))


for losses in all_team_losses:
    ax.plot(x, losses, alpha=0.3)

ax.set_xticks(x)
ax.set_xticklabels(reference_teams, rotation=90)

ax.set_xlabel('Team')
ax.set_ylabel('Average Loss (MSE)')
ax.set_title('Team-Level Average Loss Across Hyperparameter Runs')

plt.tight_layout()
plt.grid()
plt.show()

```



2.5.1 Evaluation of the team-level loss plot

The team-level loss plot indicates that the average loss remains relatively stable across different random hyperparameter searches for most teams. Although the overall variation between runs seems to be limited, certain clubs consistently exhibit higher losses or larger variance than others.

Notably, clubs such as Liverpool and Manchester City display relatively high loss values across runs. This may be an indication that these clubs have generally been among the strongest and most dominant teams in the Premier League during the analysed period.

Another notable outlier is Blackpool. Blackpool only participated in a single Premier League season during the analysed period and were relegated despite achieving 39 points, which is relatively high for a relegated team. This unusual season may contribute to unstable graph characteristics and high loss values. Furthermore, the limited amount of available match data for Blackpool likely amplifies sensitivity to small variations in the learned edge hyperparameters.

The most pronounced variance across random searches is observed for Queen Park Rangers (QPR). Between 2011 and 2015, QPR participated in only three Premier League seasons while undergoing multiple managerial changes. Such instability can significantly influence playing style and consequently in-game statistics. Since the graph edges are derived directly from statistical dominance measures, substantial tactical changes may alter the direction and magnitude of edges between seasons.

2.6 Spectral clustering on match statistics for comparison to part I clustering methods

To compare spectral clustering with the Part I methods (DBSCAN and K-means++) we take the 5-dimensional feature space where each match is represented as a point, and connect each match to its 9 nearest neighbors in feature space using the `sklearn.neighbors` module. This turns the match-statistics space into a sparse graph.

From this graph (represented by the adjacency matrix) we construct the degree matrix and the Laplacian matrix. We use the unnormalized Laplacian $L = D - A$ since it is the simplest formulation. We then compute the smallest eigenvalues and corresponding eigenvectors of the Laplacian using the eigsh function from scipy.sparse.linalg.

We use the eigenvectors associated with the smallest non-zero eigenvalues as an embedding and run K-means on that spectral representation.

This lets us compare whether the graph structure of the match statistics produces clearer clusters than DBSCAN and K-means++ in Part I.

```
[37]: X = modified_stats.values # each match is a node, features come from Part I
      ↪(modified_stats).
k = 9 # Number of nearest neighbors to connect. We use 9, since that was our
      ↪MinPts for DBSCAN.

# 1) Creates adjacency matrix using nearest neighbors. (Is sparse, so only
      ↪stores the non-zero entries --> the edges!)
A = kneighbors_graph(
    X,
    n_neighbors=k,
    mode="connectivity", # 0/1 edges
    include_self=False,
    metric="euclidean" # To use the same distance as in part I.
)

# 2) Makes the graph undirected.
A_undirected = A.maximum(A.T)

# 3) Computes the degree matrix
A = A_undirected.tocsr()
deg = np.asarray(A.sum(axis=1)).ravel()
D = sp.diags(deg, offsets=0, format="csr")

# 4) computes the laplacian matrix
L = D - A
```

```
[38]: # 5) --- We compute the smallest eigenvalues of L ---
m = 4 ## Number of eigenvalues to compute using eigsh. We want the 3 smallest
      ↪eigenvalues for clustering into 3 groups. 3 + 1 (since the first is usually
      ↪0/everything).

# eigsh solves for the eigenvalues of the L matrix. It computes the m
      ↪eigenvalues closest to sigma (0.0) --> the 10 smallest eigenvalues.
evals, evecs = sp.linalg.eigsh(L, k=m, sigma=0.0, which="LM")

# Sorts the eigenvalues in order.
order = np.argsort(evals)
```

```

evals = evals[order]
evecs = evecs[:, order]

print("Smallest eigenvalues of L:")
print(evals)

```

Smallest eigenvalues of L:
[5.71664755e-16 1.07641122e-01 2.70921755e-01 2.79144969e-01]

```

[39]: # --- does the spectral clustering using the eigenvectors --- Does kmeans on
      ↪ the eigenvectors to cluster the data into k clusters.

k_clusters = 3

# Skips the first eigenvector (corresponding to the smallest eigenvalue -->
↪ zero) and uses the next k eigenvectors as features for k-means.
U = evecs[:, 1 : k_clusters + 1] #since the first eigenvector is usually the
↪ trivial one (all ones) corresponding to the zero eigenvalue.

# Settings for k-means --> 3 clusters, fixed random seed for reproducibility.
kmeans_spec = KMeans(n_clusters=k_clusters, random_state=RANDOM_SEED,
↪ n_init="auto")

# Do k-means on the rows of U to get cluster labels.
labels_spec = kmeans_spec.fit_predict(U)

print("Spectral Clustering labels:", np.unique(labels_spec))

# Find contingency table and purity.
c_table = contingency_table(labels_spec, ground_truth, ignore_noise=False)
purity = purity_score(c_table)

print("Contingency table:\n", c_table)
print("Purity per cluster:", purity)

# Print ARI and NMI.
clustering_result3 = {
    'Spectral Graph Clustering': labels_spec,
}
for name, y_pred in clustering_result3.items():
    print(f"\n{name}:")
    ari = adjusted_rand_score(ground_truth, y_pred)
    nmi = normalized_mutual_info_score(ground_truth, y_pred)
    print(f"ARI: {ari:.4f}, NMI: {nmi:.4f}")

```

Spectral Clustering labels: [0 1 2]
Contingency table:
[[825 705 1187]

```
[ 958  473  464]
[ 585  728 1935]]
Purity per cluster: [0.43687891 0.5055409 0.59575123]
```

Spectral Graph Clustering:
ARI: 0.0608, NMI: 0.0454

```
[40]: # --- Updated for single cluster DBSCAN ---
color_map3 = {2: 'green', 1: 'blue', 0: 'red', -1: 'black'}
colors3 = [color_map3[int(val)] for val in labels]
legend_3 = {0: 'Cluster 0', -1: 'Noise'}

# --- For spectral clustering ---
color_map5 = color_map3 # reuse exactly the same colors
colors5 = [color_map5[int(val)] for val in labels_spec]
legend_5 = {0: 'Cluster 0', 1: 'Cluster 1', 2: 'Cluster 2'} # same label style
↳ as others

# --- Create figure ---
fig, ax = plt.subplots(1, 4, figsize=(16, 4))
ax = ax.flatten()

# --- Subplot 1: Match outcome ---
ax[0].scatter(embedding[:, 0], embedding[:, 1], color=colors2, alpha=0.5)
ax[0].set_xlabel("t-SNE 1")
ax[0].set_ylabel("t-SNE 2")
ax[0].set_title("t-SNE: Ground truth")
legend_elements2 = [Patch(facecolor=color_map2[key], label=legend_2[key]) for
↳ key in legend_2]
ax[0].legend(handles=legend_elements2, title="Match Outcome")

# --- Subplot 2: DBSCAN clusters ---
ax[1].scatter(embedding[:, 0], embedding[:, 1], color=colors3, alpha=0.5)
ax[1].set_xlabel("t-SNE 1")
ax[1].set_ylabel("t-SNE 2")
ax[1].set_title("t-SNE with DBSCAN Clusters")
legend_elements3 = [Patch(facecolor=color_map3[key], label=legend_3[key]) for
↳ key in legend_3]
ax[1].legend(handles=legend_elements3, title="DBSCAN Cluster")

# --- Subplot 3: K-means++ clusters ---
ax[2].scatter(embedding[:, 0], embedding[:, 1], color=colors4, alpha=0.5)
ax[2].set_xlabel("t-SNE 1")
ax[2].set_ylabel("t-SNE 2")
ax[2].set_title("t-SNE with K-means++ Clusters")
legend_elements4 = [Patch(facecolor=color_map4[key], label=legend_4[key]) for
↳ key in legend_4]
```


```

ax[2].legend(handles=legend_elements4, title="K-means++ Cluster")

# --- Subplot 4: Spectral clusters ---
ax[3].scatter(embedding[:, 0], embedding[:, 1], color=colors5, alpha=0.5)
ax[3].set_xlabel("t-SNE 1")
ax[3].set_ylabel("t-SNE 2")
ax[3].set_title("t-SNE with Spectral Clusters")
legend_elements5 = [Patch(facecolor=color_map5[key], label=legend_5[key]) for
    ↪key in legend_5]
ax[3].legend(handles=legend_elements5, title="Spectral Cluster")

plt.tight_layout()
plt.show()

```



K-means++ and spectral clustering perform similarly, while DBSCAN still performs worse when evaluated against the ground truth. Spectral clustering finds a different graph-based partition and gives slightly higher ARI, but K-means++ gives slightly higher purity and NMI. However the conclusion from Part I (clustering) remains. The clusters do not align well with the ground truth.

2.7 Further investigations could include

Comparing title winning team games, instead of all games across the season. This could reveal: - The unusual Leicester title win in 2015/2016, or - The 2018/2019 season were was widely considered to be unlucky not winning the title.

2.8 How the feedback was from checkpoint 1 was addressed

- A manhattan distance was introduced when deciding the hyperparameters for DBSCAN. However, due to the sensitivity in the number of clusters when changing ϵ , this was quickly disregarded.
- ARI and NMI evaluation of both DBSCAN and K-Means++ was implemented to introduce some quantitative analysis of the perfnace of algorithms.

3 Pattern mining (Part III)

Pattern mining objective:

The objective of the pattern mining analysis is to explore whether association rule mining can uncover meaningful relationships between bookmakers' odds across different seasons. By discretizing the odds features into categorical bins, we aim to identify frequent itemsets and generate association rules that reveal how certain odds patterns co-occur across matches. This analysis can provide insights into the consistency of bookmakers' odds and potentially highlight any systematic biases or trends in their offerings over time.

3.1 Data loading and preprocessing and feature engineering for pattern mining

The idea is to do a discretization of the odds features to see if we can mine association rules between the bookmakers. Note that there is a slight difference in the odds features across the seasons, as some seasons have odds from more bookmakers than others. To address this we choose a subset of seasons where the same bookmakers are represented.

```
[41]: def load_common_data(folder_path, min_files=10, encoding="latin1",
    ↪ print_progress=False):
    """
    Loads and concatenates CSV files that share a common set of columns.

    The function scans all CSV files in the specified folder, identifies their
    available columns, and greedily selects files such that the intersection of
    shared columns is maximized while including at least `min_files` files.
    Only the common columns are loaded and concatenated into a single dataframe.

    Parameters:
        folder_path (str): Path to the directory containing the CSV files.
        min_files (int): Minimum number of compatible CSV files to include in
    ↪ the final dataset.
        encoding (str): File encoding used when reading the CSV files. Default
    ↪ is "latin1".
        print_progress (bool): If True, prints information during execution.

    Returns:
        final_df (pd.DataFrame): Concatenated dataframe containing only the
    ↪ shared columns across the
        selected files, with an additional "Season" column indicating the
    ↪ source file.
        selected_files (list[str]): List of filenames included in the final
    ↪ merged dataset.
        current_common (list[str]): List of column names shared across all
    ↪ selected files.
    """
    file_columns = {}

    # --- Step 1: Read headers safely ---
    for file in os.listdir(folder_path):
        if file.endswith(".csv"):
```

```

    path = os.path.join(folder_path, file)
    try:
        df = pd.read_csv(path, nrows=0, encoding=encoding)
        file_columns[file] = set(df.columns)
    except Exception as e:
        print(f"Skipping {file} during scan: {e}")

    if len(file_columns) < min_files:
        raise ValueError(f"Not enough readable files. Found {len(file_columns)}")

    # --- Step 2: Start with file that has the most columns ---
    sorted_files = sorted(file_columns.items(), key=lambda x: len(x[1]),
        reverse=True)

    selected_files = [sorted_files[0][0]]
    current_common = set(sorted_files[0][1])

    if print_progress:
        print(f"\n Starting with: {os.path.splitext(selected_files[0])[0]} {len(current_common)} cols")

    remaining_files = {f: cols for f, cols in sorted_files[1:]}

    # --- Step 3: Greedily add files that maximize remaining common columns ---
    while len(selected_files) < min_files and remaining_files:
        best_file = None
        best_intersection = set()

        for f, cols in remaining_files.items():
            intersection = current_common.intersection(cols)

            if len(intersection) > len(best_intersection):
                best_file = f
                best_intersection = intersection

        if best_file is None or len(best_intersection) == 0:
            print("No more compatible files can be added without losing all columns.")
            break

        selected_files.append(best_file)
        current_common = best_intersection

        if print_progress:
            print(f" + Added: {os.path.splitext(best_file)[0]} {len(current_common)} common cols")

```

```

        del remaining_files[best_file]

# --- Final report ---
if print_progress:
    print(f"\n Final selected {len(selected_files)} files:")
for f in selected_files:
    if print_progress:
        print(f" - {os.path.splitext(f)[0]} ({len(file_columns[f])} cols)")

if print_progress:
    print(f"\n Final shared columns: {len(current_common)}")
    print(sorted(current_common))

# --- Step 4: Load data ---
dataframes = []

for file in selected_files:
    path = os.path.join(folder_path, file)
    try:
        df = pd.read_csv(path, usecols=lambda c: c in current_common,
↪encoding=encoding)

        season = os.path.splitext(file)[0]
        df["Season"] = season


        dataframes.append(df)
    except Exception as e:
        print(f"Skipping {file} during load: {e}")

final_df = pd.concat(dataframes, ignore_index=True)

return final_df, selected_files, list(current_common)

final_df, num_cols = [], []
for i in range(1, 24):
    df, _, cols = load_common_data(dataPath, min_files=i)
    final_df.append(df)
    num_cols.append(len(cols))
plt.plot(range(1, 24), num_cols, marker='o')
plt.xlabel('Number of Files Included')
plt.ylabel('Number of Common Columns')
plt.title('Common Columns vs. Number of Files Included')
plt.grid()
plt.show()

```



3.1.1 Multi-Season Dataset Construction (Greedy Alignment)

A greedy approach was used to build a consistent dataset across multiple football seasons.

The process starts with the season containing the most columns, then iteratively adds seasons that preserve the largest possible overlap of shared columns. At each step, the set of common columns is updated to ensure all selected seasons remain compatible.

After selecting the final group of seasons, only the columns shared across all of them are kept, and each season is loaded into a single combined DataFrame.

This ensures a consistent feature set across seasons while maximizing the number of usable datasets included.

We have chosen to include 7 seasons in our pattern mining analysis, which provides a good balance between data richness and consistency.

```
[42]: data, selected_files, common_cols = load_common_data(dataPath, min_files=7)
print(f"Selected files: {sorted(selected_files)}")
print(f"Number of common columns: {len(common_cols)}")
print(f"Number of seasons loaded: {data['Season'].nunique()}")
```



```

Selected files: ['20192020premierleague.csv', '20202021premierleague.csv',
'20212022premierleague.csv', '20222023premierleague.csv',
'20232024premierleague.csv', '20242025premierleague.csv',
'20252026premierleague.csv']
Number of common columns: 87
Number of seasons loaded: 7

```

```

[43]: list_of_bookmakers = ['1XB', 'B365', 'BF', 'BFD', 'BMGM', 'BV', 'BS', 'BW',
                             'CL', 'GB', 'IW', 'LB', 'PS', 'SO', 'SB', 'SJ', 'SY',
                             'VC', 'WH']

cols = common_cols
present_bookmakers = set()

for b in list_of_bookmakers:
    if any(col.startswith(b) for col in cols):
        present_bookmakers.add(b)

print("Present bookmakers:", sorted(present_bookmakers))
print("Count:", len(present_bookmakers))

```

```

Present bookmakers: ['B365', 'BW', 'PS']
Count: 3

```

Here we define valid columns as those that start with a known bookmaker prefix and ends with a result related suffix (H, D, A). That is we are only concerned with betting odds columns for home win, away win, and draw. This is due to the fact that including other odds columns (as odds for over/under 2,5 goals) is gonna make the interesting patterns drown in a sea of irrelevant or trivial patterns.

Furthermore we chose to include the referee column, as we want to see if there are any patterns between the referees and the odds. However this would be very surprising.

```

[44]: prefixes = ["B365", "BW", "PS"]
suffixes = ["A", "H", "D"]
valid_cols = {p + s for p in prefixes for s in suffixes}
relevant_cols = [c for c in common_cols if c in valid_cols] + ['Referee']
print(f"Relevant columns include: {relevant_cols}\nnumber of columns:
    ↳{len(relevant_cols)}")

```

```

Relevant columns include: ['PSH', 'PSA', 'BWH', 'PSD', 'BWA', 'B365A', 'BWD',
'B365H', 'B365D', 'Referee']
number of columns: 10

```

3.1.2 Feature engineering and discretization

In order to do pattern mining on the odds features, we need to discretize them. Therefore we needed to investigate the distribution of the odds features to find a good way to discretize them.

```
[45]: fig, ax = plt.subplots(3, 3, figsize=(20, 10))
ax = ax.flatten()


for i, c in enumerate(sorted(relevant_cols)):
    if i >= len(ax): # safety in case more columns than plots
        break

    data_col = data[c].dropna()

    ax[i].hist(data_col, bins=100, edgecolor='white', color='k')
    ax[i].set_title(f"{c}")
    ax[i].set_xlabel("Odds")
    ax[i].set_ylabel("Frequency")
    # -- log scale for better visibility of low-frequency bins --
    #ax[i].set_yscale('log')
    ax[i].set_xlim(0, 30)
    ax[i].grid(True)

# hide unused subplots (if any)
for j in range(i + 1, len(ax)):
    ax[j].axis("off")

plt.tight_layout()
plt.show()
```



Based on the distributions above, it does not make sense to use a single discretization scheme for all columns, as *draw* odds differ substantially from *home* and *away* odds. We therefore construct separate bins per outcome type (H, D, A).

Instead of binning each column independently, we pool odds across all bookmakers for each outcome and compute quantile-based bins. This ensures consistent bins across bookmakers while capturing the distinct distributional shape of each outcome type.

```
[46]: def compute_quantile_bins(data, relevant_cols, suffixes=["H", "D", "A"],
    ↪n_bins=10):
    """
    Compute quantile-based bin edges for each outcome type (H, D, A),
    pooling odds across all bookmakers for that outcome type.
    Parameters:
        data (pd.DataFrame): DataFrame containing the odds columns
        relevant_cols (list): List of column names to consider for pooling
        suffixes (list): List of outcome suffixes
        n_bins (int): Number of quantile bins to create (default 5)

    Returns:
        A dict like: {'H': (bins, labels), 'D': (bins, labels), 'A': (bins,
    ↪labels)}
    """
    bin_specs = {}

    for suffix in suffixes:
        # Pool all odds for this outcome across all bookmakers
        cols_for_suffix = [c for c in relevant_cols if c.endswith(suffix)]
        pooled = pd.concat([data[c].dropna() for c in cols_for_suffix])

        # Compute quantile edges (0%, 20%, 40%, 60%, 80%, 100%)
        quantiles = np.linspace(0, 1, n_bins + 1)
        edges = pooled.quantile(quantiles).values

        # Ensure strictly increasing edges (can have ties at extremes)
        edges = np.unique(edges)
        edges[0] = 0 # open lower bound
        edges[-1] = np.inf # open upper bound

        # Labels like "1.42-1.75", "1.75-2.10", etc.
        labels = [
            f"{edges[i]:.2f}-{edges[i+1]:.2f}".replace("inf", "∞")
            for i in range(len(edges) - 1)
        ]

        bin_specs[suffix] = (edges, labels)
        print(f"[{suffix}] edges: {np.round(edges, 3)} → {len(labels)} bins")

    return bin_specs
```

```

def discretize_odds_quantile(series, bin_specs):
    """
    Discretize a single odds column using precomputed quantile bin specs.
    Detects outcome type from the column name's last character.
    Parameters:
        series (pd.Series): pandas Series of odds to discretize
        bin_specs (dict): dict of {suffix: (edges, labels)} for discretization
    Returns:
        pd.Series: pandas Series of categorical bin labels corresponding to the
    ↪ input odds
    """
    suffix = series.name[-1] # 'H', 'D', or 'A'
    bins, labels = bin_specs[suffix]
    return pd.cut(series, bins=bins, labels=labels)

# Compute shared bins per outcome type, pooled across bookmakers
bin_specs = compute_quantile_bins(data, relevant_cols, n_bins=10)

# Apply to each column
discrete_data = data.copy()
for c in relevant_cols:
    if c[-1] in bin_specs:
        discrete_data[c] = discretize_odds_quantile(discrete_data[c], bin_specs)

```

```

[H] edges: [0.    1.34 1.54 1.75 2.    2.26 2.6   3.    3.75 5.25  inf] → 10 bins
[D] edges: [0.    3.3   3.41 3.54 3.66 3.8   4.    4.33 4.75 5.774  inf]
→ 10 bins
[A] edges: [0.    1.63 2.    2.37 2.75 3.2   3.75 4.6   6.    8.75  inf] → 10 bins

```

Computed bins is shown above. One for each outcome type (H, D, A), with bin edges determined by quantiles of the pooled odds data.

```

[47]: for suffix in ["H", "D", "A"]:
    # get all columns for this outcome
    cols_for_suffix = [c for c in relevant_cols if c.endswith(suffix)]
    for sample_col in cols_for_suffix:
        print(f"\n{sample_col} bin distribution (%):")
        print(
            (discrete_data[sample_col].value_counts(normalize=True) * 100)
            .sort_index()
            .round(1)
        )

```

PSH bin distribution (%):

PSH

0.00-1.34	9.7
1.34-1.54	10.1
1.54-1.75	9.6

1.75-2.00	9.5
2.00-2.26	9.5
2.26-2.60	11.0
2.60-3.00	9.3
3.00-3.75	10.8
3.75-5.25	9.9
5.25-∞	10.4

Name: proportion, dtype: float64

BWH bin distribution (%):

BWH

0.00-1.34	10.1
1.34-1.54	9.9
1.54-1.75	10.6
1.75-2.00	11.2
2.00-2.26	9.0
2.26-2.60	10.6
2.60-3.00	10.1
3.00-3.75	9.5
3.75-5.25	9.8
5.25-∞	9.3

Name: proportion, dtype: float64

B365H bin distribution (%):

B365H

0.00-1.34	10.5
1.34-1.54	10.0
1.54-1.75	11.2
1.75-2.00	10.1
2.00-2.26	9.3
2.26-2.60	10.2
2.60-3.00	10.2
3.00-3.75	9.5
3.75-5.25	9.6
5.25-∞	9.4

Name: proportion, dtype: float64

PSD bin distribution (%):

PSD

0.00-3.30	6.9
3.30-3.41	7.8
3.41-3.54	10.8
3.54-3.66	10.6
3.66-3.80	9.8
3.80-4.00	11.2
4.00-4.33	11.8
4.33-4.75	9.0
4.75-5.77	10.8

5.77-∞ 11.5
Name: proportion, dtype: float64

BWD bin distribution (%):

BWD

0.00-3.30	14.3
3.30-3.41	8.9
3.41-3.54	9.3
3.54-3.66	9.6
3.66-3.80	12.6
3.80-4.00	9.6
4.00-4.33	8.6
4.33-4.75	8.5
4.75-5.77	9.6
5.77-∞	9.1

Name: proportion, dtype: float64

B365D bin distribution (%):

B365D

0.00-3.30	14.7
3.30-3.41	7.6
3.41-3.54	10.2
3.54-3.66	9.6
3.66-3.80	14.0
3.80-4.00	8.4
4.00-4.33	9.3
4.33-4.75	8.2
4.75-5.77	8.6
5.77-∞	9.4

Name: proportion, dtype: float64

PSA bin distribution (%):

PSA

0.00-1.63	9.7
1.63-2.00	9.8
2.00-2.37	9.6
2.37-2.75	9.8
2.75-3.20	10.2
3.20-3.75	9.9
3.75-4.60	10.1
4.60-6.00	9.9
6.00-8.75	10.5
8.75-∞	10.5

Name: proportion, dtype: float64

BWA bin distribution (%):

BWA

0.00-1.63	10.3
-----------	------

1.63-2.00	11.0
2.00-2.37	9.7
2.37-2.75	10.7
2.75-3.20	10.1
3.20-3.75	10.0
3.75-4.60	9.5
4.60-6.00	10.6
6.00-8.75	8.9
8.75-∞	9.5

Name: proportion, dtype: float64

B365A bin distribution (%):

B365A	
0.00-1.63	10.6
1.63-2.00	10.8
2.00-2.37	9.2
2.37-2.75	10.5
2.75-3.20	9.4
3.20-3.75	10.2
3.75-4.60	9.7
4.60-6.00	10.9
6.00-8.75	8.9
8.75-∞	9.8

Name: proportion, dtype: float64

Although we used quantile-based bins pooled across all bookmakers, individual bookmakers' odds distributions are not identical. Note the percentage of data in each bin to differ across bookmakers, even though the bins themselves represent global quantiles. This is a natural consequence of skewed distributions between bookmakers.

3.1.3 Applying the the Apriori algorithm

The Apriori algorithm is applied to the discretized odds data to identify frequent itemsets and generate association rules. The algorithm iteratively builds itemsets of increasing size, starting from single items (individual odds categories) and expanding to larger combinations. At each step, it prunes itemsets that do not meet a specified minimum support threshold. Once frequent itemsets are identified, association rules are generated. Here we defined a function *is_relevant* to filter the generated rules based on certain criteria to filter out irrelevant or trivial rules.

```
[48]: def make_transactions(df, cols):
    """
    Converts the discretized odds columns into a list of transactions for
    ↪association rule mining.

    Each transaction is a list of items like "B365H=0-1.34" representing the
    ↪bin for that column in that row.

    Parameters:
```

```

    df (pd.DataFrame): The dataframe containing the discretized odds_
↳ columns.
    cols (list): List of column names to include in the transactions.
Returns:
    list of lists: Each inner list is a transaction containing items
    """
    T = []

    for _, row in df[cols].iterrows():
        items = []

        for c in cols:
            val = row[c]
            if pd.notna(val):
                items.append(f"{c}={val}")

        T.append(items)

    return T
T = make_transactions(discrete_data, relevant_cols)

rule = lambda lhs, rhs: "%s => %s" % (str(lhs), str(rhs)) # For generating_
↳ rule strings

def compute_support(Ck: list, T: list, still_applicable=None):
    """
    Function that computes the support of each candidate itemset in Ck with_
↳ respect to the transactions in T.
    Parameters:
        Ck (list): list of candidate itemsets of size k
        T (list): list of transactions (each transaction is a list of items)
        still_applicable (list): a list of boolean values indicating whether_
↳ each transaction is applicable
    Returns:
        counts (dict): a dictionary of counts corresponding to the support of_
↳ each candidate itemset in Ck
    """
    if still_applicable is None:
        still_applicable = [True] * len(T)

    counts = {} # Dictionary to store counts of each candidate itemset slide 40
    for row, t in enumerate(T):
        if not still_applicable[row]: continue
        found_any = False
        for c in Ck:
            if set.issubset(set(c), t):
                if c in counts:

```



```

        counts[c] += 1
    else:
        counts[c] = 1
        found_any = True
    still_applicable[row] = found_any
    return counts

def extend_prefix_tree(Ck_prev):
    """
    Function that extends the prefix tree by generating new candidate itemsets
    of size k+1 from the previous candidate itemsets of size k.
    Parameters:
        Ck_prev (list): list of candidate itemsets of size k
    Returns:
        Ck (list): list of candidate itemsets of size k+1
    """
    Ck = set()
    #join step
    for itemset in Ck_prev:
        its1 = tuple(sorted(itemset))
        for itemset2 in Ck_prev:
            its2 = tuple(sorted(itemset2))
            if its1[:-1] == its2[:-1] and its1 < its2: Ck.add(its1 + its2[-1:])
    # Join step: if the first k-1 items are the same and the last item is
    # different, we can join them
    #prune step
    to_remove = set()
    for c in Ck:
        for subset in combinations(c, len(c) - 1):
            if subset not in Ck_prev:
                to_remove.add(c)
                break
    for c in to_remove:
        Ck.remove(c)
    return Ck

def powerset(iterable):
    """
    Function that generates the powerset of a given iterable.
    Parameters:
        iterable (iterable (e.g., list, set, etc.)): for which to generate
    the powerset
    Returns:
        A list of all subsets of the input iterable
    """
    s = list(iterable)

```

```

    out = chain.from_iterable(combinations(s, r) for r in range(1, len(s))) #
    ↪Generate all non-empty subsets
    return out

def is_relevant(A, B):
    def parse(item):
        col, val = item.split("=")
        bookmaker = col[:-1]
        outcome = col[-1]
        value = val.replace(" ", "")
        return bookmaker, outcome, value

    A_parsed = [parse(x) for x in A]
    B_parsed = [parse(x) for x in B]

    all_parsed = A_parsed + B_parsed

    # 1. Reject if rule involves all three outcomes (A, H, D)
    outcomes = {o for _, o, _ in all_parsed}
    if outcomes == {"A", "H", "D"}:
        return False

    # 2. Remove same-market duplication across bookmakers
    A_market = {(o, v) for _, o, v in A_parsed}
    B_market = {(o, v) for _, o, v in B_parsed}
    if A_market & B_market:
        return False

    # 3. Reject "mirrored" home/away rules
    #If any outcome in A is 'H' and in B is 'A' (or vice versa), reject
    A_outcomes = {o for _, o, _ in A_parsed}
    B_outcomes = {o for _, o, _ in B_parsed}
    if ("H" in A_outcomes and "A" in B_outcomes) or ("A" in A_outcomes and "H"
    ↪in B_outcomes):
        return False

    return True

def apriori_algorithm(T, support=0.05, min_confidence=0.7):
    """
        Apriori algorithm for mining association rules.
        Parameters:
            T:                A list of lists, each inner list will contain
            ↪integer-item-ids.                Example: T = [[1, 2, 5], [2, 3, 4], [1, 6]]

            support:          The proportion of occurrences needed to keep
            ↪itemsets.
    """

```

```

        min_confidence: Minimum confidence for the algorithm to output the
    ↪ rule.

    Returns:
        rules ([[str, float]]): A list of tuples (rule_str, confidence)
    """
    rules = []
    # --- Generate frequent itemsets using the Apriori algorithm ---
    N = len(T) # Total number of transactions
    # Construct the initial candidate itemsets of size 1
    items = set(i for t in T for i in t) # Get all unique items in the
    ↪ transactions
    C1 = set((i,) for i in items) # Convert to set of tuples for consistency
    ↪ with later itemsets
    counts = compute_support(C1, T) # Compute support for itemsets of size 1
    L1 = [c for c in counts if counts[c] / N >= support] # Filter itemsets by
    ↪ support
    L = [L1]
    support_dict = {c: counts[c] / N for c in L1} # Store support values for
    ↪ itemsets of size 1
    # Second step
    k = 2
    while L[k-2]: # While there are itemsets of size k-1
        Ck = extend_prefix_tree(L[k-2]) # Generate candidate itemsets of size k
        if not Ck: break # If no candidates, stop

        counts = compute_support(Ck, T) # Compute support for itemsets of size k
        Lk = []
        for i in Ck:
            supp = counts.get(i, 0) / N # Get support for itemset i, default to
            ↪ 0 if not found
            if supp >= support:
                Lk.append(i) # Keep itemsets that meet the support threshold
                support_dict[i] = supp # Store support value for itemset i
            if not Lk: break # If no itemsets meet the support threshold, stop
        L.append(Lk) # Add the frequent itemsets of size k to the list
        k += 1
    # --- Generate rules from the frequent itemsets ---
    for Lk in L:
        for itemset in Lk:
            if len(itemset) < 2: continue # We need at least 2 items to form a
            ↪ rule

            for A in powerset(itemset):
                B = tuple(sorted(set(itemset) - set(A))) # Get the remaining
                ↪ items as B
                if not B: continue # Skip if B is empty

```

```

        if not is_relevant(A, B):
            continue

        conf = support_dict[itemset] / support_dict[A] # Calculate
↳confidence of the rule A => B
        if conf >= min_confidence:
            rules.append((rule(A, B), conf)) # Store the rule along
↳with its confidence
        return rules

def pretty_print_rules(rules):
    """
    Pretty-print association rules as a table.
    Parameters:
        rules ([[str, float]]): list of tuples (rule_str, confidence)
    Returns:
        None (displays a sorted DataFrame of rules)
    """
    data = []
    for r, conf in rules:
        # Parse rule string of the form "('A',) => ('B',)"
        lhs, rhs = r.split(" => ")
        lhs = lhs.strip("(").replace("'", "").replace(",", "")
        rhs = rhs.strip("(").replace("'", "").replace(",", "")
        data.append({"LHS": lhs, "RHS": rhs, "Confidence": round(conf, 3)})

    df = pd.DataFrame(data)
    display(df.sort_values("Confidence", ascending=False).
↳reset_index(drop=True))

support = 0.08
min_confidence = 0.70
rules = apriori_algorithm(T, support=support, min_confidence=min_confidence)
print(f"Generated {len(rules)} rules with support >= {support} and confidence
↳>= {min_confidence}")
print("Sample rules:")
pretty_print_rules(rules)

```

Generated 20 rules with support >= 0.08 and confidence >= 0.7

Sample rules:

	LHS \
0	B365H=0.00-1.34 BWH=0.00-1.34 PSH=0.00-1.34
1	BWH=0.00-1.34 PSH=0.00-1.34
2	B365H=0.00-1.34 PSH=0.00-1.34
3	PSH=0.00-1.34
4	B365H=0.00-1.34 BWH=0.00-1.34

5	B365A=8.75- ∞ PSA=8.75- ∞
6	BWH=0.00-1.34
7	B365H=0.00-1.34
8	B365A=8.75- ∞
9	PSA=8.75- ∞
10	PSD=5.77- ∞
11	PSD=5.77- ∞
12	PSD=5.77- ∞
13	PSD=5.77- ∞
14	PSD=5.77- ∞
15	PSD=5.77- ∞
16	PSD=5.77- ∞
17	PSD=5.77- ∞
18	PSD=5.77- ∞
19	PSD=5.77- ∞

	RHS	Confidence
0	PSD=5.77- ∞	0.932
1	PSD=5.77- ∞	0.932
2	PSD=5.77- ∞	0.925
3	PSD=5.77- ∞	0.921
4	PSD=5.77- ∞	0.889
5	PSD=5.77- ∞	0.873
6	PSD=5.77- ∞	0.871
7	PSD=5.77- ∞	0.846
8	PSD=5.77- ∞	0.839
9	PSD=5.77- ∞	0.824
10	B365H=0.00-1.34	0.787
11	PSH=0.00-1.34	0.780
12	B365H=0.00-1.34 PSH=0.00-1.34	0.780
13	PSA=8.75- ∞	0.752
14	BWH=0.00-1.34	0.734
15	B365H=0.00-1.34 BWH=0.00-1.34	0.731
16	B365A=8.75- ∞	0.727
17	B365A=8.75- ∞ PSA=8.75- ∞	0.724
18	BWH=0.00-1.34 PSH=0.00-1.34	0.724
19	B365H=0.00-1.34 BWH=0.00-1.34 PSH=0.00-1.34	0.724

3.2 Evaluation of the generated rules

Generally it is hard to find a good definition of a non trivial rule in this context. We made a function that filters out trivial rules based on the following criteria: - **Reject rules involving all three outcomes (A, H, D):**

If a rule includes “Away win (A)”, “Home win (H)”, and “Draw (D)” simultaneously, it is rejected. This avoids rules like:

(‘B365A=8.75- ∞ ’, ‘B365H=0.00-1.34’) => (‘PSD=5.77- ∞ ’,)

This rule is trivial because it simply states that if the odds for an away win are very high and

the odds for a home win are very low, then the odds for a draw are also very high, which is not insightful.

- **Remove duplication across bookmakers:**

If any (outcome, value) pair appears in both side A and side B, the rule is rejected. This rule prevents trivial rules that simply restate the same condition across different bookmakers, such as:

$((\text{'BWA}=8.75-\infty'),) \Rightarrow (\text{'PSA}=8.75-\infty'), 0.98).$

This rule is trivial because it just says that if one bookmaker has high odds for an away win, then another bookmaker also has high odds for an away win, which is not insightful.

- **Reject home/away rules:**

If side A contains “H” and side B contains “A”, or vice versa, the rule is rejected. This prevents trivial rules that simply mirror home and away conditions, such as:

$((\text{'BWH}=5.25-\infty'),) \Rightarrow (\text{'PSA}=0.00-1.63'), 0.96).$

This rule is trivial because it just states that if the odds for a home win are very high, then the odds for an away win are very low, which is not insightful.

With these rejections in place and by tuning the parameters support and min_confidence, we succeeded in filtering out enough rules to be able to manually inspect the remaining rules for interesting patterns. We found some interesting patterns, such as:

$((\text{'PSD}=5.77-\infty'),) \Rightarrow (\text{'PSA}=8.75-\infty'), 0.75)$

This rule suggests that when the draw is considered very unlikely by Pinnacle, the odds for an away win are very high (75% of the time). By extension, this implies that a home win becomes more likely under the same conditions. This pattern indicates that Pinnacle may systematically favor home teams in its odds. This pattern can also be found directly on the form:

$((\text{'PSD}=5.77-\infty'),) \Rightarrow (\text{'PSH}=0.00-1.34'), 0.78)$

and with different bookmakers involved:

$((\text{'PSD}=5.77-\infty'),) \Rightarrow (\text{'BWH}=0.00-1.34', \text{'PSH}=0.00-1.34'), 0.724)$

3.2.1 How the feedback from checkpoint 2 was addressed

In the feedback we were told to change the graph direction in the pagerank part due to risk of data leakage. This concern was not real, since it was only a matter of bad phrasing in a markdown cell. The following was done to clarify: - Restating the text in the markdown cell “More sophisticated graph-based method”, to clarify that the direction of the edge is not determined by the result rather by the sign of the dot product between the feature vector and the hyperparameters. The “random_search” function takes as input a boolean “data_leakage”. When “data_leakage” is set to false, the zeroth and seventh entry in the hyperparameter vector is hardcoded to zero before searching for best_params. These two entries correspond to “FTG_diff” and “AvgOdds_diff” which are highly related to the actual outcome of the match. - Adding a comment in the crucial if statement in the function “random_search” where the hardcoding is happening. - Adding a comment in the function “match_row_to_edge” where the direction is determined

Regarding the suggestion of introducing graph spectral clustering: In the bottom of part II, we tried to make graph spectral clustering, but the results didn't seem much better than when we used K-Means++ and DBSCAN. So afterall the conclusion from Part I remains.

3.2.2 Further Investigations

Since the dataset contains a large number of matches, it could be interesting to investigate whether a team's current form influences the bookmakers' odds. However, to perform this analysis, the following steps are required:

- Aggregate the most recent *prev_match* matches for each team into a single current form rating, computed at each match date.
- Normalise each team's odds relative to a team strength ranking, for instance using *PageRank*. This allows us to isolate changes in odds over time that are driven by form rather than opponent strength — without this normalisation, odds would primarily reflect the quality of the opponent rather than the team's recent performance.